




Class #14: Convolutional Neural Networks

Machine Learning (CS 135)


Neural Networks for Images

- ▶ A regular feed forward network can sometimes prove problematic for image-processing tasks
 - ▶ Given a (100×100) pixel color image, each with 3 color-channel (e.g. RGB) values, we end up with many, many weights to be learned
 - ▶ In addition, a 1-D weight-vector doesn't carry any real information about **spatial relationships** between image features (edges, blocks of color, ...)



Convolutional Neural Networks (CNNs)

- ▶ To capture image dynamics, and expand what the networks can do, we organize neurons into stacks of 3-dimensional volumes
 - ▶ Each is connected to later volumes, filtering and flattening down to the usual final ($C \times 1$) classification-output layer (where C is the number of classes)




Types of Layers in CNNs

- ▶ **INPUT:** as in a typical NN, each neuron corresponds to a single **input feature-value**
 - ▶ Only the 3-D arrangement is different
- ▶ **OUTPUT:** again, as in a typical NN, these are **fully-connected** layers
 - ▶ Each neuron is connected to all of those in the volume above
 - ▶ Each computes a function, like the sigmoid (*softmax*), typically giving probabilities for each of the possible output classes
- ▶ **OTHER:** layers between can play different possible roles
 1. **CONVOLUTION:** transformations on sub-regions
 2. **RELU:** application of the $\max(0, x)$ function
 3. **POOLING:** down-sampling to reduce volume size

Deep Convolutional Networks

- ▶ For complex image-classification tasks, we may use many layers, combining the types in varying orders



Convolutional (CONV) Layers


- ▶ The core innovation in a CNN is the idea of a **spatial filter**, which is a 3-D volume where:
 1. Each neuron in one layer computes a function on a proper **sub-region** of the layer above
 2. We form the CONV layer by “tiling” the prior layer, in (possibly) overlapping sub-regions
 3. Every neuron in one layer shares a **single set** of weights, and so computes the same function

- ▶ Two main decisions in building such a layer:
 1. What **size** of sub-region should we use?
 2. What is our **stride**; i.e., **how far** do we move over each time we connect our next sub-region?

Result of filter function


Convolutional Layer

5 x 5 pixel filter



Input: (28 x 28)


Suppose we choose a sub-region
size of (5 x 5) pixels



Result of filter function


Convolutional Layer

Stride: move
2 pixels right




Input: (28 x 28)

Suppose we also choose a
stride-value = 2



Convolutional Layer:
(14 x 14)

“Off-edge” pixel
values all set to 0



Input: (28 x 28)

Since stride = 2, the result is a layer with half
as many neurons in each dimension




A Full Convolutional Layer


The 3-dimensional CONV layer consists of a stack of N such filters, of dimensionality:
 $(14 \times 14 \times N)$

Every neuron in each filter-layer **shares** a single set of common weights, applied to inputs, with the products summed as usual.

N different convolutions

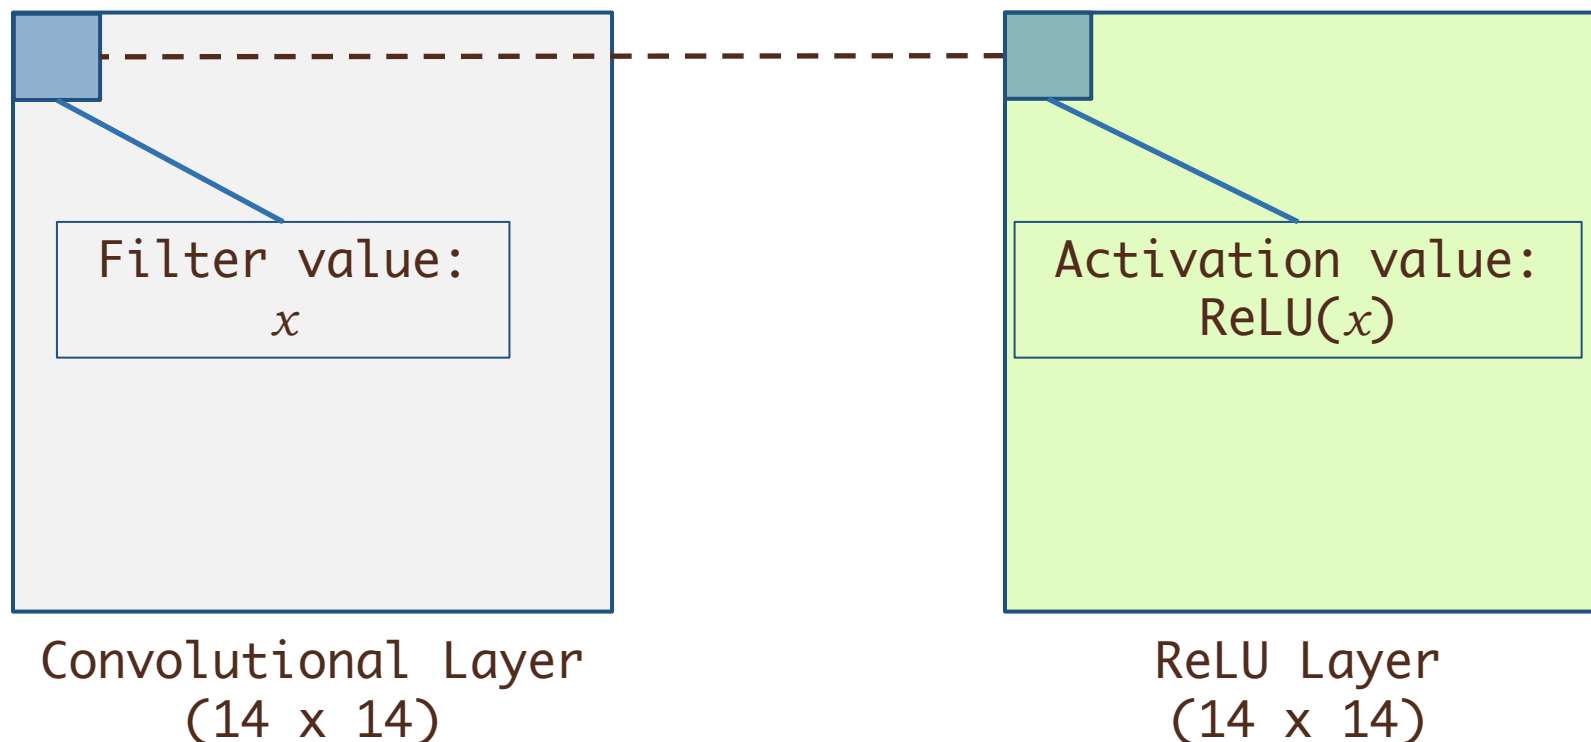


(28×28)




ReLU (Activation) Layers

- ▶ CONV layer may or may not change input size (depends upon stride)
- ▶ ReLU layer keeps size the same, simply applying its function to neurons
 - ▶ ReLU is very popular, but other activation function layers are allowed




Combining Layers

Using a 3-dimensional convolutional layer of multiple filters means that we will have a matching number of activation layers.




Pooling Layers

- ▶ While CONV and ReLU layers can compute more complex functions, POOL layers **down-sample** a region, reducing it to something simpler (usually its MAX value)




Pooling Layers

- ▶ Again, we stride across the layer, reducing the overall size by avoiding overlap
- ▶ Most common approach: (2×2) region, with stride = 2



Combining Layers

Again, each layer is 3-dimensional (until the final output layer).



Uses of CNNs and Other Deep Networks

- ▶ Convolutional networks have become increasingly popular for image and other spatial data
- ▶ Browser-based demos:
<https://cs.stanford.edu/people/karpathy/convnetjs/>
- ▶ A variety of applications of neural network models to a number of research problems
<https://youtu.be/Bui3DWs02h4>
<https://youtu.be/hPKJBXkyTKM>
<https://youtu.be/aKSILzbAqJs>
- ▶ Cat drawings!
<https://affinelayer.com/pixsrv/>