

Thu Apr 13 11:02 PM

DataSpell-EAP File Edit Code Cell Run Kernel Environment Tools VCS Window Help

project1 Version control

notebook.ipynb implementation.py

Add Configuration... Managed: http://localhost:8890 Python 3 (ipykernel) Trusted

10 28

Project 1 – Support Vector Machine Classification

NAME(S): Morgan Rockett, Prithvi Shahani, Anthony Wong

DATE: 4/03/23

What will we do?

Using gradient descent, we will build a Support Vector Machine to find the optimal hyperplane that maximizes the margin between two toy data classes.

What are some use cases for SVMs?

-Classification, regression (time series prediction, etc.), outlier detection, clustering

How does an SVM compare to other ML algorithms?

Classifiers: (a) Logistic Regression, (b) SVM, and (c) Multi-Layer Perception (MLP)

- As a rule of thumb, SVMs are great for relatively small data sets with fewer outliers.
- Other algorithms (Random forests, deep neural networks, etc.) require more data but almost always develop robust models.
- The decision of which classifier to use depends on your dataset and the general complexity of the problem.
- "Premature optimization is the root of all evil (or at least most of it) in programming." – Donald Knuth, CS Professor (Turing award speech 1974)

Other Examples

- Learning to use Scikit-learn's SVM function to classify images https://github.com/ksopyla/svm_mnist_digit_classification
- Pulse classification, more useful dataset <https://github.com/akasantony/pulse-classification-svm>

```
## What is a Support Vector Machine?
```

It's a supervised machine learning algorithm that can be used for both classification and regression problems. But it's usually used for classification. Given two or more labeled data classes, it acts as a discriminative classifier, formally defined by an optimal hyperplane that separates all the classes. New examples mapped into that space can then be categorized based on which side of the gap they fall.

project1 > notebook.ipynb

30:1 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

Managed: http://localhost:8890 Python 3 (ipykernel) Trusted

What is a Support Vector Machine?

It's a supervised machine learning algorithm that can be used for both classification and regression problems. But it's usually used for classification. Given two or more labeled data classes, it acts as a discriminative classifier, formally defined by an optimal hyperplane that separates all the classes. New examples mapped into that space can then be categorized based on which side of the gap they fall.

What are Support Vectors?

The diagram shows two scatter plots. The left plot, labeled 'Small Margin', shows a narrow gap between two classes of points (green and blue) separated by a vertical decision boundary. The right plot, labeled 'Large Margin', shows a wider gap between the same two classes separated by a diagonal decision boundary. The points closest to the boundary are labeled 'Support Vectors'.

Support vectors are the data points nearest to the hyperplane, the points of a data set that, if removed, would alter the position of the dividing hyperplane. Because of this, they can be considered the critical elements of a data set; they help us build our SVM.

What is a hyperplane?

Hyperplanes as decision surfaces

- A hyperplane is a linear decision surface that splits the space into two parts;
- It is obvious that a hyperplane is a binary classifier.

The left diagram shows a 2D grid with red and green data points. A blue line (hyperplane) separates them. The right diagram shows a 3D grid with red and green data points. A blue plane (hyperplane) separates them. Below these, a caption states: "A hyperplane in \mathbb{R}^n is an $n-1$ dimensional subspace".

Geometry tells us that a hyperplane is a subspace of one dimension less than its ambient space. For instance, a hyperplane of an n -dimensional space is a flat subset with size $n - 1$. By its nature, it separates the space in half.

Linear vs nonlinear classification?

project1 > notebook.ipynb

30:1 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

Managed: http://localhost:8890 Python 3 (ipykernel) Trusted

Linear vs nonlinear classification?

Sometimes our data is linearly separable. That means for N classes with M features. We can learn a mapping that is a linear combination. (like $y = mx + b$). Or even a multidimensional hyperplane ($y = x + z + b + q$). No matter how many dimensions/features a set of classes have, we can represent the mapping using a linear function.

But sometimes it is not. Like if there was a quadratic mapping. Luckily for us, SVMs can efficiently perform a non-linear classification using what is called the kernel trick.

More on this as a Bonus question comes at the end of notebook.

All right, let's get to the building!

Instructions

In this assignment, you will implement a support vector machine (SVM) from scratch, and you will use your implementation for multiclass classification on the MNIST dataset.

In `implementation.py` implement the SVM class. In the `fit` function, use `scipy.optimize.minimize` (see documentation) to solve the constrained optimization problem:

$$\begin{aligned} \text{maximize}_a \quad & \sum_{i=1}^n a_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n a_i a_j y_i y_j (x_i \cdot x_j) \\ \text{subject to} \quad & a_i \geq 0, i = 1, \dots, n \\ & \sum_{i=1}^n a_i y_i = 0 \end{aligned}$$

Note: An SVM is a convex optimization problem. Using to solve the equation above will be computationally expensive given larger datasets. CS 168 Convex Optimization is a course to take later if interested in optimization and the mathematics and intuition that drives it.

```
In 26 1 import numpy as np
2 import pandas as pd
3
4 from scipy.io import loadmat
5 from implementation import SVM, linear_kernel, nonlinear_kernel
6 # from solution import SVM, linear_kernel, nonlinear_kernel
7 from sklearn.datasets import make_blobs
8 from sklearn.svm import SVC
9 from sklearn.metrics import accuracy_score
10 from sklearn.linear_model import LogisticRegression
```

project1 > notebook.ipynb

30:1 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

In 26

```
1 import numpy as np
2 import pandas as pd
3
4 from scipy.io import loadmat
5 from implementation import SVM, linear_kernel, nonlinear_kernel
6 # from solution import SVM, linear_kernel, nonlinear_kernel
7 from sklearn.datasets import make_blobs
8 from sklearn.svm import SVC
9 from sklearn.metrics import accuracy_score
10 from sklearn.linear_model import LogisticRegression
11 from sklearn.metrics import confusion_matrix
12
13 import matplotlib
14 import matplotlib.pyplot as plt
15
16 %load_ext autoreload
17 %autoreload 2
Executed in 250ms, 13 Apr at 22:24:13
```

The autoreload extension is already loaded. To reload it, use:
%reload_ext autoreload

Step 1 – Get Data

In 27

```
1 # Input data - of the form [Bias term, x_1 value, x_2 value]
2 X = np.array([
3     [1, -2, 4],
4     [1, 4, 1],
5     [1, 1, 6],
6     [1, 2, 4],
7     [1, 6, 2],
8 ])
9
10 # Associated output labels - first 2 examples are labeled '-1' and last 3 are labeled '+1'
11 y = np.array([-1,-1,1,1,1])
12
13 # Let's plot these examples on a 2D graph!
14 # Plot the negative samples (the first 2)
15 plt.scatter(X[:,1][y== -1], X[:,2][y== -1], s=120, marker='.', linewidths=2)
16 # Plot the positive samples (the last 3)
17 plt.scatter(X[:,1][y== 1], X[:,2][y== 1], s=120, marker='+', linewidths=2)
18
19 # Print a possible hyperplane, that is separating the two classes.
20 # we'll two points and draw the line between them (naive guess)
21 plt.plot([-2,6],[6,0.5])
22 plt.xlabel("x1")
23 plt.ylabel("x2")
```

project1 > notebook.ipynb

30:1 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

In 27

```
1 # Input data - of the form [Bias term, x_1 value, x_2 value]
2 X = np.array([
3     [1, -2, 4],
4     [1, 4, 1],
5     [1, 1, 6],
6     [1, 2, 4],
7     [1, 6, 2],
8 ])
9
10 # Associated output labels - first 2 examples are labeled '-1' and last 3 are labeled '+1'
11 y = np.array([-1,-1,1,1])
12
13 # Let's plot these examples on a 2D graph!
14 # Plot the negative samples (the first 2)
15 plt.scatter(X[:,1][y== -1], X[:,2][y== -1], s=120, marker='_', linewidths=2)
16 # Plot the positive samples (the last 3)
17 plt.scatter(X[:,1][y== 1], X[:,2][y== 1], s=120, marker='+', linewidths=2)
18
19 # Print a possible hyperplane, that is separating the two classes.
20 # we'll two points and draw the line between them (naive guess)
21 plt.plot([-2,6],[6,0.5])
22 plt.xlabel(r"$x_1$")
23 plt.ylabel(r"$x_2$")
24 plt.show()
```

Executed in 243ms, 13 Apr at 22:24:13

project1 > notebook.ipynb

30:1 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

SVM basics

SVM using scikit-learn.

```
In 28 1 result = SVC(kernel="Linear")
2 result.fit(X, y.ravel())
3
4 print("scikit-learn indices of support vectors:", result.support_)
5
Executed in 144ms, 13 Apr at 22:24:13
```

scikit-learn indices of support vectors: [0 1 3 4]

Implement and test SVM to sklearn's version (20 points)

Compare the indices of support vectors from scikit-learn with `implementation.py` using toy data.

```
In 29 1 # TODO: implement SVM, along with linear_kernel
2
3 result = SVC(kernel="linear")
4 result.fit(X, y)
5
6 print("scikit-learn indices of support vectors:", result.support_)
7
8 svm = SVM(kernel=linear_kernel)
9 svm.fit(X, y)
Executed in 124ms, 13 Apr at 22:24:13
```

scikit-learn indices of support vectors: [0 1 3 4]

message: Optimization terminated successfully

success: True

status: 0

fun: -0.62499999898889

x: [1.475e-01 4.774e-01 1.468e-15 4.088e-01 2.162e-01]

nit: 6

jac: [-4.000e+00 -4.000e+00 5.500e+00 4.000e+00 4.000e+00]

nfev: 39

njev: 6

[1.47549954e-01 4.77448320e-01 1.46822658e-15 4.08827517e-01
2.16170757e-01]
[-1.44328993e-15 4.99986204e-01 1.00000345e+00]
-3.99982745681246

Out 29 <implementation.SVM at 0x296503e50>

project1 > notebook.ipynb

30:1 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

In 30 1 print("implementation.py indices of support vectors:",\n2 np.array(range(y.shape[0]))[svm.a>1e-8])\n3\n4 if (result.support_ != np.array(range(y.shape[0]))[svm.a>1e-8]).all():\n5 raise Exception("The calculation is wrong")

Executed in 101ms, 13 Apr at 22:24:13

implementation.py indices of support vectors: [0 1 3 4]

Compare the weights assigned to the features from scikit-learn with implementation.py.

In 31 1 #1000 - other sections were done for you, specify the variables to print, find the difference, and check it is within reasonable error from that of sklearn's version.\n2 print("scikit-learn weights assigned to the features:", result.coef_)\n3 print("implementation.py weights assigned to the features:", svm.w)\n4\n5 diff = abs(result.coef_ - svm.w)\n6 if (diff > 1e-3).any():\n7 raise Exception("The calculation is wrong")

Executed in 81ms, 13 Apr at 22:24:13

scikit-learn weights assigned to the features: [[0. 0.5 0.99969451]]
implementation.py weights assigned to the features: [-1.44328993e-15 4.99986204e-01 1.00000345e+00]

Compare the bias weight from scikit-learn with implementation.py.

In 32 1 print("scikit-learn bias weight:", result.intercept_)\n2 print("implementation.py bias weight:", svm.b)\n3\n4 diff = abs(result.intercept_ - svm.b)\n5 if (diff > 1e-3).all():\n6 raise Exception("The calculation is wrong")

Executed in 62ms, 13 Apr at 22:24:13

scikit-learn bias weight: [-3.99915989]
implementation.py bias weight: -3.999982745681246

Compare the predictions from scikit-learn with implementation.py.

In 33 1 X_test = np.array([\n2 [4, 4, -1],\n3 [1, 3, -1],\n4])

project1 > notebook.ipynb

149:1 LF UTF-8 4 spaces ml_0135

The screenshot shows the DataSpell-EAP IDE interface. At the top, the menu bar includes: File, Edit, Code, Cell, Run, Kernel, Environment, Tools, VCS, Window, Help. The status bar at the bottom right shows: 149:1 LF UTF-8 4 spaces ml_0135. The main workspace displays a Jupyter notebook titled "notebook.ipynb".

In 33:

```
1 X_test = np.array([
2     [4, 4, -1],
3     [1, 3, -1]
4 ])
5 print("scikit-learn predictions:", result.predict(X_test))
6 print("implementation.py predictions:", svm.predict(X_test))
7
8 if (svm.predict(X_test) != result.predict(X_test)).all():
9     raise Exception("The calculation is wrong")
Executed in 44ms, 13 Apr at 22:24:13
```

Output:

```
scikit-learn predictions: [-1 -1]
implementation.py predictions: [-1. -1.]
```

Using SKLearn's SVM (one-versus-the-rest)

You can load the data with `scipy.io.loadmat`, which will return a Python dictionary containing the test and train data and labels.

In 34:

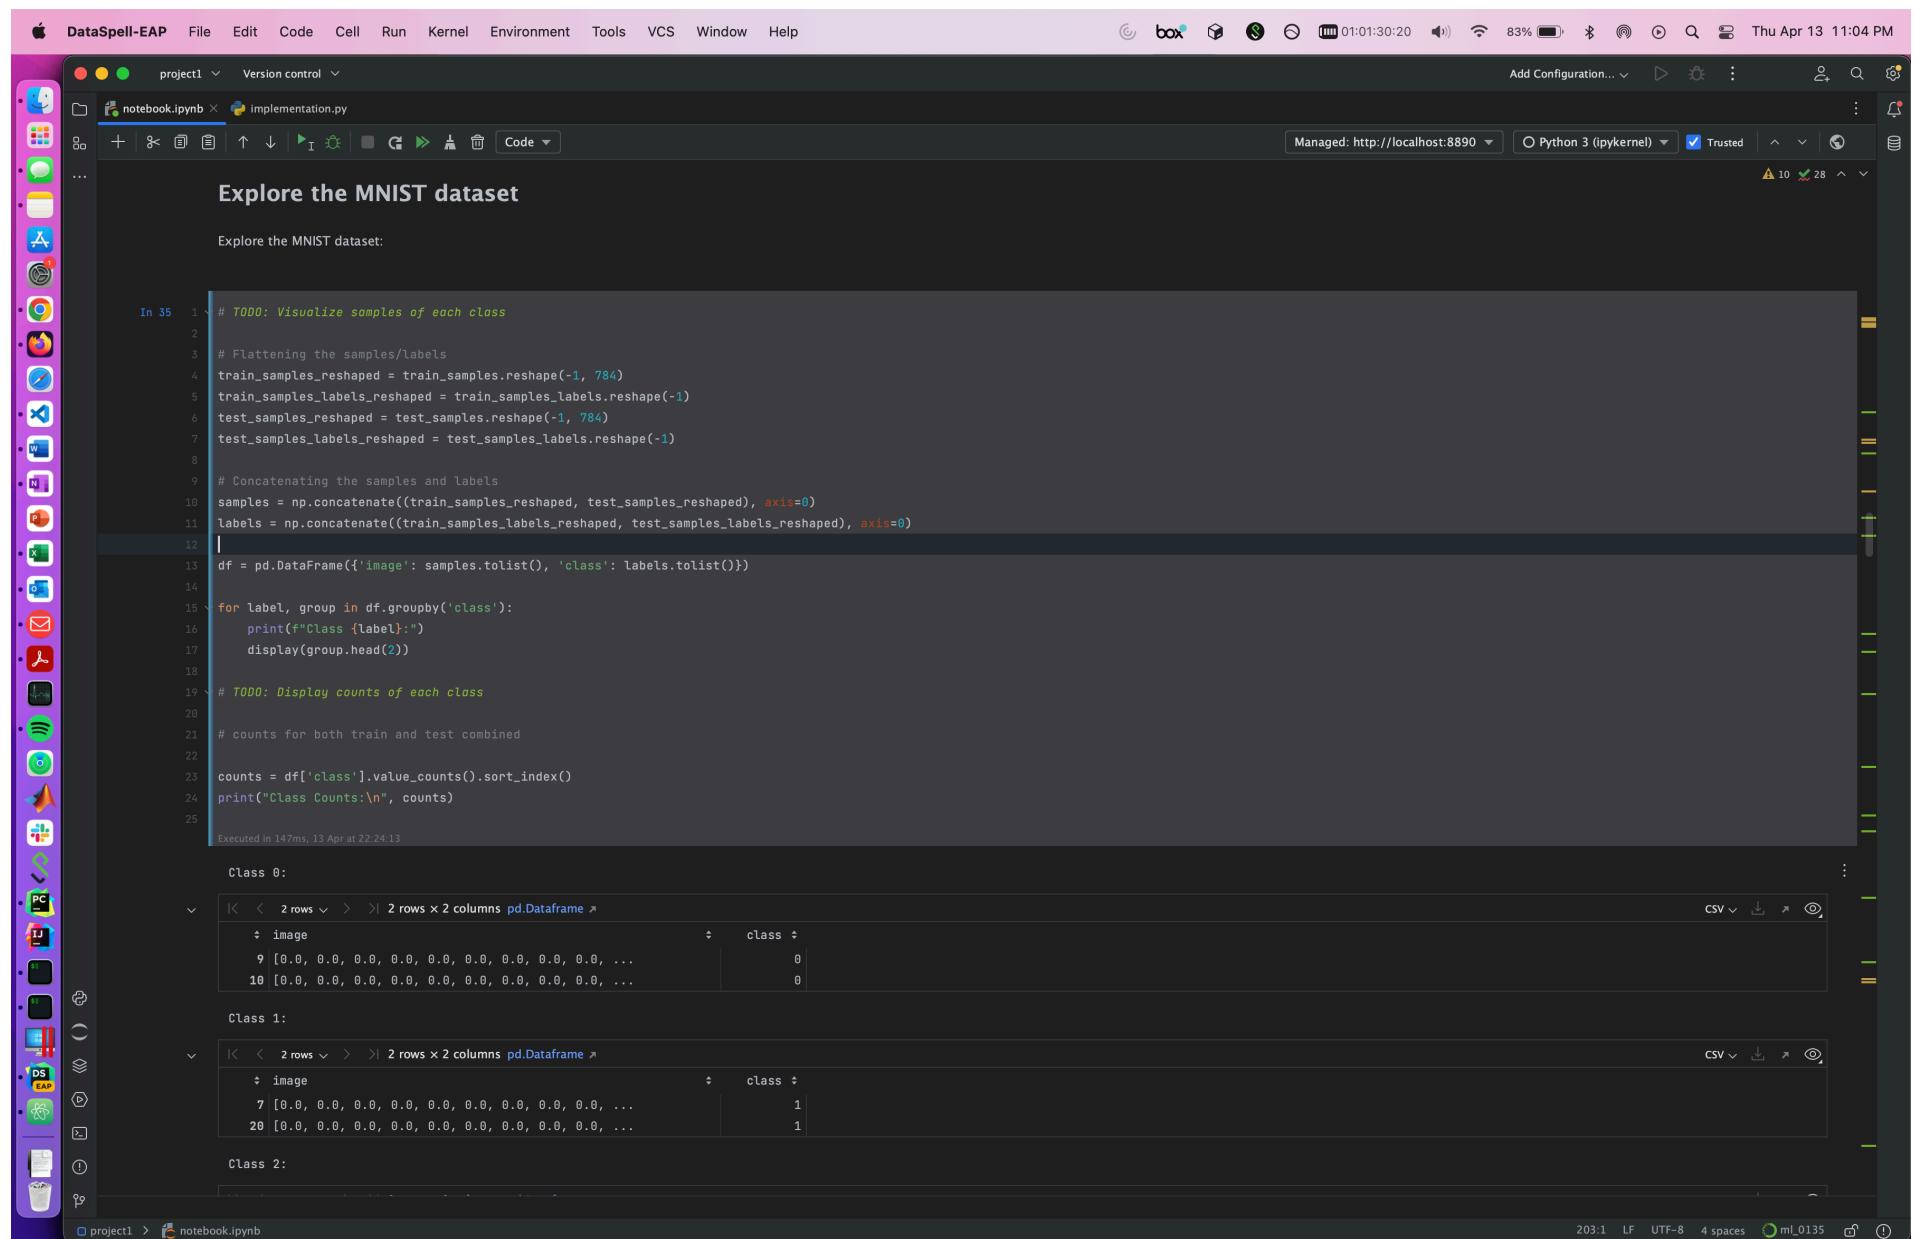
```
1 mnist = loadmat('data/MNIST.mat')
2 train_samples = mnist['train_samples']
3 train_samples_labels = mnist['train_samples_labels']
4 test_samples = mnist['test_samples']
5 test_samples_labels = mnist['test_samples_labels']
Executed in 139ms, 13 Apr at 22:24:13
```

Explore the MNIST dataset

Explore the MNIST dataset:

In 35:

```
1 # TODO: Visualize samples of each class
2
3 # Flattening the samples/labels
4 train_samples_reshaped = train_samples.reshape(-1, 784)
5 train_samples_labels_reshaped = train_samples_labels.reshape(-1)
6 test_samples_reshaped = test_samples.reshape(-1, 784)
7 test_samples_labels_reshaped = test_samples_labels.reshape(-1)
8
9 # Concatenating the samples and labels
10 samples = np.concatenate((train_samples_reshaped, test_samples_reshaped), axis=0)
11 labels = np.concatenate((train_samples_labels_reshaped, test_samples_labels_reshaped), axis=0)
12
13
14 df = pd.DataFrame({'image': samples.tolist(), 'class': labels.tolist()})
15
```



project1 > notebook.ipynb

File Edit Code Cell Run Kernel Environment Tools VCS Window Help

Managed: http://localhost:8890 Python 3 (ipykernel) Trusted

Thu Apr 13 11:04 PM

Class 2:

	image	class
8	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	2
11	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	2

Class 3:

	image	class
14	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	3
19	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	3

Class 4:

	image	class
15	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	4
29	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	4

Class 5:

	image	class
12	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	5
28	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	5

Class 6:

	image	class
3	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	6
4	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	6

Class 7:

	image	class
1	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	7
2	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	7

Class 8:

	image	class
13	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	8
17	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	8

203:1 LF UTF-8 4 spaces ml_0135

A screenshot of the DataSpell-EAP IDE interface. The main window displays eight separate dataframes, each representing a different class (Class 2 through Class 8). Each dataframe has two columns: 'image' (containing lists of numerical values) and 'class' (containing integer values). The dataframes are presented in a vertical stack, each preceded by a section header. The interface includes a top navigation bar with tabs like File, Edit, Code, Cell, Run, Kernel, Environment, Tools, VCS, Window, and Help. A status bar at the bottom shows file statistics like '203:1 LF UTF-8 4 spaces' and a file name 'ml_0135'. On the left, there's a vertical toolbar with various icons and a sidebar showing project files.

project1 Version control

notebook.ipynb implementation.py

Code

Class 9:

	image	class
0	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	9
5	[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...]	9

Class Counts:

0	468
1	573
2	523
3	528
4	525
5	439
6	481
7	511
8	459
9	493

Name: class, dtype: int64

one-versus-the-rest (15 Points) and analysis

Using your implementation, compare multiclass classification performance of *one-versus-the-rest*

Create your own implementation of *one-versus-the-rest* and *one-versus-one*. Do not use sklearns multiclass SVM.

```
In 36 1 # TODO loop over classes training one_vs_the_rest()
2
3 classes = np.unique(train_samples_labels)
4 train_probs_svm = np.zeros((len(classes), len(train_samples)))
5 test_probs_svm = np.zeros((len(classes), len(test_samples)))
6 X_train = train_samples.copy()
7 X_test = test_samples.copy()
8
9 for i, curr_class in enumerate(classes):
10     y_train = np.where(train_samples_labels == curr_class, 1, -1)
11     svm = SVC(kernel="linear", probability=True)
12     svm.fit(train_samples, y_train.ravel())
13     y_test = np.where(test_samples_labels == curr_class, 1, -1)
14     train_probs_svm[i] = svm.predict_proba(train_samples)[:, 1]
15     test_probs_svm[i] = svm.predict_proba(test_samples)[:, 1]
16     print("{}-versus-the-rest, train accuracy: {}".format(i, svm.score(X_train, y_train)))
17     print("{}-versus-the-rest, test accuracy: {}".format(i, svm.score(X_test, y_test)))
18
19 # TODO save all the prediction probability by predict_prob() for the following function
20 # Hint: svm = SVC(kernel="linear", probability=True)
```

project1 > notebook.ipynb

203:1 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

Code

Managed: http://localhost:8890 Python 3 (ipykernel) Trusted

10 28

```
0-versus-the-rest, train accuracy: 0.99475
0-versus-the-rest, test accuracy: 0.985
1-versus-the-rest, train accuracy: 0.9985
1-versus-the-rest, test accuracy: 0.994
2-versus-the-rest, train accuracy: 0.97825
2-versus-the-rest, test accuracy: 0.97
3-versus-the-rest, train accuracy: 0.9745
3-versus-the-rest, test accuracy: 0.962
4-versus-the-rest, train accuracy: 0.9815
4-versus-the-rest, test accuracy: 0.977
5-versus-the-rest, train accuracy: 0.97425
5-versus-the-rest, test accuracy: 0.966
6-versus-the-rest, train accuracy: 0.98875
6-versus-the-rest, test accuracy: 0.98
7-versus-the-rest, train accuracy: 0.985
7-versus-the-rest, test accuracy: 0.971
8-versus-the-rest, train accuracy: 0.9635
8-versus-the-rest, test accuracy: 0.951
9-versus-the-rest, train accuracy: 0.965
9-versus-the-rest, test accuracy: 0.961
```

Determine the accuracy

In 37

```
1 predicted_training_labels = np.argmax(train_probs_svm, axis=0)
2 predicted_test_labels = np.argmax(test_probs_svm, axis=0)
3
4 train_accuracy = accuracy_score(train_samples.labels, predicted_training_labels)
5 test_accuracy = accuracy_score(test_samples.labels, predicted_test_labels)
6 print("Train accuracy: {:.2f}".format(100*train_accuracy))
7 print("Test accuracy: {:.2f}".format(100*test_accuracy))
Executed in 18ms, 13 Apr at 22:25:23
```

Train accuracy: 91.80
Test accuracy: 88.60

The parameter $C > 0$ controls the tradeoff between the size of the margin and the slack variable penalty. It is analogous to the inverse of a regularization coefficient. Include in your report a brief discussion of how you found an appropriate value.

In 38

```
1 # Hint: Try using np.logspace for hyperparameter tuning
2 # TODO: Find an appropriate value of C.
3 train_accuracies = list()
4 test_accuracies = list()
5
6 for i in np.logspace(-4, 3, 10):
7     train_probs = np.empty((0, len(train_samples)))
```

project1 > notebook.ipynb

242:53 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

The parameter $C > 0$ controls the tradeoff between the size of the margin and the slack variable penalty. It is analogous to the inverse of a regularization coefficient. Include in your report a brief discussion of how you found an appropriate value.

```
In 38 1 # Hint: Try using np.logspace for hyperparameter tuning
2 # TODO: Find an appropriate value of C.
3 train_accuracies = list()
4 test_accuracies = list()
5
6 for i in np.logspace(-4, 3, 10):
7     train_probs = np.empty((0, len(train_samples)))
8     test_probs = np.empty((0, len(test_samples)))
9
10    for j in range(len(classes)):
11        indices_isclass = np.where(train_samples_labels == classes[j])[0]
12        indices_notclass = np.where(train_samples_labels != classes[j])[0]
13
14        y_train_isclass = np.zeros(len(train_samples_labels))
15        y_train_isclass[indices_isclass] = 1
16        y_train_isclass[indices_notclass] = 0
17
18        svm = SVC(C = i, kernel="linear", probability=True)
19        svm.fit(train_samples, y_train_isclass)
20        pred_train_prob = svm.predict_proba(train_samples)[:,1]
21        pred_test_prob = svm.predict_proba(test_samples)[:,1]
22
23        train_probs = np.append(train_probs, [pred_train_prob], axis=0)
24        test_probs = np.append(test_probs, [pred_test_prob], axis=0)
25
26        predicted_training_labels_hat = np.argmax(train_probs, axis=0)
27        predicted_test_labels_hat = np.argmax(test_probs, axis=0)
28        train_accuracy_hat = accuracy_score(train_samples_labels, predicted_training_labels_hat)
29        test_accuracy_hat = accuracy_score(test_samples_labels, predicted_test_labels_hat)
30
31        train_accuracies.append(train_accuracy_hat)
32        test_accuracies.append(test_accuracy_hat)
33
34    # print(train_accuracies)
35    # print(test_accuracies)
36    for i in range(10):
37
38        print("{}-versus-the-rest, train accuracy: {}".format(i, train_accuracies[i]))
39        print("{}-versus-the-rest, test accuracy: {}".format(i, test_accuracies[i]))

```

Executed in 12m, 13 Apr at 22:38:08

0-versus-the-rest, train accuracy: 0.84725
0-versus-the-rest, test accuracy: 0.812
1-versus-the-rest, train accuracy: 0.86925
1-versus-the-rest, test accuracy: 0.839
2-versus-the-rest, train accuracy: 0.8685

project1 > notebook.ipynb

242:53 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

Code

```
0-versus-the-rest, train accuracy: 0.84725
0-versus-the-rest, test accuracy: 0.812
1-versus-the-rest, train accuracy: 0.86925
1-versus-the-rest, test accuracy: 0.839
2-versus-the-rest, train accuracy: 0.8685
2-versus-the-rest, test accuracy: 0.842
3-versus-the-rest, train accuracy: 0.86875
3-versus-the-rest, test accuracy: 0.841
4-versus-the-rest, train accuracy: 0.88125
4-versus-the-rest, test accuracy: 0.853
5-versus-the-rest, train accuracy: 0.9115
5-versus-the-rest, test accuracy: 0.881
6-versus-the-rest, train accuracy: 0.938
6-versus-the-rest, test accuracy: 0.894
7-versus-the-rest, train accuracy: 0.965
7-versus-the-rest, test accuracy: 0.881
8-versus-the-rest, train accuracy: 0.98525
8-versus-the-rest, test accuracy: 0.841
9-versus-the-rest, train accuracy: 0.99725
9-versus-the-rest, test accuracy: 0.824
```

Provide details on how you found an appropriate value.

We are examining the C value from a range between 10^{-4} to 10^3 . This iterates through 10 classes to predict the highest probability by using the argmax. Then, we see the accuracies of these methods from the accuracy_score function. These metrics are stored in train and test accuracies and eventually the code shows a comparison of the accuracies.

Plot accuracies for train and test using logspace for x-axis (i.e., C values)

```
In 39 1 # TODO: Plot the result.
2
3 import matplotlib.pyplot as plt
4
5 C_values = np.logspace(-4, 3, 10)
6
7 plt.plot(C_values, train_accuracies, label="Accuracy on train")
8 plt.plot(C_values, test_accuracies, label="Accuracy on test")
9
10 plt.gca().axes.set_xscale("log")
11
12 plt.axvline(x=1, color='green', linestyle='--')
for i in range(10)
```

project1 > notebook.ipynb

Managed: http://localhost:8890 Python 3 (ipykernel) Trusted

10 28

294:81 LF UTF-8 4 spaces ml_0135

DataSpell-EAP

File Edit Code Cell Run Kernel Environment Tools VCS Window Help

project1 Version control

notebook.ipynb implementation.py

Add Configuration... <img

project1 Version control

notebook.ipynb implementation.py

Code

What does this graph tell us about the importance of our C value?

TODO: Analyze the plot above:

This graph clearly shows us the importance of the C value with regards to controlling the maximizing margin and minimizing classification error trade-off. As we can see, the higher C is, the narrower the margin is, but this could lead to overfitting. On the other hand, the smaller C is, the wider the margin, and thus less sensitive to outliers and could lead to underfitting. As observed, from 10^0 onward, the accuracy of the train and test data diverges.

(10 Points)

In addition to calculating percent accuracy, generate multiclass confusion matrices as part of your analysis.

```
In 40 1 train_predictions = list()
2 test_predictions = list()
3 # TODO
4 from sklearn.metrics import confusion_matrix
5
6 conf_matrices = list()
7
8 for i in np.logspace(-4, 3, 10):
9     train_probs = np.empty((0, len(train_samples)))
10    test_probs = np.empty((0, len(test_samples)))
11
12    for j in range(len(classes)):
13        indices_isclass = np.where(train_samples_labels == classes[j])[0]
14        indices_notclass = np.where(train_samples_labels != classes[j])[0]
15
16        y_train_isclass = np.zeros(len(train_samples_labels))
17        y_train_isclass[indices_isclass] = 1
18        y_train_isclass[indices_notclass] = 0
19
20        svm = SVC(C = i, kernel="linear", probability=True)
21        svm.fit(train_samples, y_train_isclass)
22        pred_train_prob = svm.predict_proba(train_samples)[:, 1]
23        pred_test_prob = svm.predict_proba(test_samples)[:, 1]
24
25        train_probs = np.append(train_probs, [pred_train_prob], axis=0)
26        test_probs = np.append(test_probs, [pred_test_prob], axis=0)
27
28    predicted_training_labels_hat = np.argmax(train_probs, axis=0)
29    predicted_test_labels_hat = np.argmax(test_probs, axis=0)
30    train_accuracy_hat = accuracy_score(train_samples_labels, predicted_training_labels_hat)
31    test_accuracy_hat = accuracy_score(test_samples_labels, predicted_test_labels_hat)
32
33 for i in range(10)
```

project1 > notebook.ipynb

294:81 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

Code

Managed: http://localhost:8890 Python 3 (ipykernel) Trusted

10 28

```
predicted_training_labels_hat = np.argmax(train_probs, axis=0)
predicted_test_labels_hat = np.argmax(test_probs, axis=0)
train_accuracy_hat = accuracy_score(train_samples_labels, predicted_training_labels_hat)
test_accuracy_hat = accuracy_score(test_samples_labels, predicted_test_labels_hat)
conf_matrix_hat = confusion_matrix(test_samples_labels, predicted_test_labels_hat)

train_accuracies.append(train_accuracy_hat)
test_accuracies.append(test_accuracy_hat)
conf_matrices.append(conf_matrix_hat)

print(train_accuracies)
print(test_accuracies)
for i, matrix in enumerate(conf_matrices):
    print("Confusion matrix for C={}: \n{}".format(np.logspace(-4, 3, 10)[i], matrix))

```

Executed in 12m, 13 Apr at 22:50:51

[0.84725, 0.86925, 0.8685, 0.86875, 0.88125, 0.9115, 0.938, 0.965, 0.98525, 0.99725, 0.848, 0.86925, 0.8685, 0.8685, 0.88125, 0.91175, 0.93825, 0.966, 0.985, 0.99725]

[0.812, 0.839, 0.842, 0.841, 0.853, 0.881, 0.894, 0.881, 0.841, 0.824, 0.815, 0.839, 0.842, 0.84, 0.852, 0.881, 0.894, 0.88, 0.837, 0.823]

Confusion matrix for C=0.0001:

[[79 0 0 0 0 2 0 4 1]
[0 117 0 1 0 0 0 0 0]
[0 3 93 5 1 0 0 3 6 2]
[0 0 1 95 0 6 3 4 5 1]
[1 1 1 0 92 1 3 0 4 5]
[4 0 1 23 2 40 2 1 16 3]
[4 0 1 0 2 3 74 0 3 0]
[1 9 3 2 3 0 0 78 0 3]
[0 0 2 6 3 0 0 2 70 3]
[0 1 0 3 2 0 0 4 5 77]]

Confusion matrix for C=0.0005994842503189409:

[[80 0 0 0 0 4 0 2 0]
[0 118 0 1 1 1 0 0 1 0]
[1 2 94 2 1 1 0 4 6 2]]

Evaluation (15 points)

Now we will report our results and compare to other algorithms. Usually compare with a handful Logistic regression

Create your own implementation of *one-versus-the-rest* and *one-versus-one*. Do not use sklearns multiclass Logistic Regression.

```
In 41 1 # TODO
2 from sklearn.linear_model import LogisticRegression
3
4 classes = np.unique(train_samples_labels)
5 train_probs_lr = np.zeros((len(classes), len(train_samples)))
6 test_probs_lr = np.zeros((len(classes), len(test_samples)))
7
8 for i in range(10)
```

project1 > notebook.ipynb

294:81 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

Code

Managed: http://localhost:8890 Python 3 (ipykernel) Trusted

10 28

Confusion matrix for C=1000.0:

79	0	0	1	0	3	1	0	0	2
0	120	0	0	0	1	0	0	1	0
0	5	91	5	0	0	1	3	5	3
0	0	4	87	1	13	1	4	0	5
1	1	2	0	87	0	3	0	6	8
3	0	1	8	2	68	1	0	8	1
3	1	3	0	1	1	74	2	1	1
2	2	4	3	0	0	0	82	1	5
0	1	1	4	6	5	0	3	63	3
0	2	2	2	7	0	0	4	3	72

Evaluation (15 points)

Now we will report our results and compare to other algorithms. Usually compare with a handful Logistic regression

Create your own implementation of *one-versus-the-rest* and *one-versus-one*. Do not use sklearns multiclass Logistic Regression.

```
In 41 1 # TODO
2 from sklearn.linear_model import LogisticRegression
3
4 classes = np.unique(train_samples_labels)
5 train_probs_lr = np.zeros((len(classes), len(train_samples)))
6 test_probs_lr = np.zeros((len(classes), len(test_samples)))
7 X_train = train_samples.copy()
8 X_test = test_samples.copy()
9
10 for i, curr_class in enumerate(classes):
11     y_train = np.where(train_samples_labels == curr_class, 1, -1)
12     svm = LogisticRegression()
13     svm.fit(train_samples, y_train.ravel())
14     y_test = np.where(test_samples_labels == curr_class, 1, -1)
15     train_probs_lr[i] = svm.predict_proba(train_samples)[:, 1]
16     test_probs_lr[i] = svm.predict_proba(test_samples)[:, 1]
17     print("{}-versus-the-rest, train accuracy: {}".format(i, svm.score(X_train, y_train)))
18     print("{}-versus-the-rest, test accuracy: {}".format(i, svm.score(X_test, y_test)))
19
Executed in 2s, 13 Apr at 22:50:53
```

1-versus-the-rest, test accuracy: 0.993
2-versus-the-rest, train accuracy: 0.96925
2-versus-the-rest, test accuracy: 0.967
3-versus-the-rest, train accuracy: 0.968
3-versus-the-rest, test accuracy: 0.959
4-versus-the-rest, train accuracy: 0.977
4-versus-the-rest, test accuracy: 0.971
5-versus-the-rest, train accuracy: 0.95925

for i in range(10)

project1 > notebook.ipynb

294:81 LF UTF-8 4 spaces ml_0135

project1 Version control

notebook.ipynb implementation.py

Code

Executed in 2s, 13 Apr at 22:50:53

```
0-versus-the-rest, train accuracy: 0.99125
0-versus-the-rest, test accuracy: 0.985
1-versus-the-rest, train accuracy: 0.98725
1-versus-the-rest, test accuracy: 0.993
2-versus-the-rest, train accuracy: 0.96925
2-versus-the-rest, test accuracy: 0.967
3-versus-the-rest, train accuracy: 0.968
3-versus-the-rest, test accuracy: 0.959
4-versus-the-rest, train accuracy: 0.977
4-versus-the-rest, test accuracy: 0.971
5-versus-the-rest, train accuracy: 0.95925
5-versus-the-rest, test accuracy: 0.945
6-versus-the-rest, train accuracy: 0.9815
6-versus-the-rest, test accuracy: 0.976
7-versus-the-rest, train accuracy: 0.98025
7-versus-the-rest, test accuracy: 0.969
8-versus-the-rest, train accuracy: 0.9525
8-versus-the-rest, test accuracy: 0.946
9-versus-the-rest, train accuracy: 0.95475
9-versus-the-rest, test accuracy: 0.961
```

Create a table comparing model accuracy on train and test data.

Method	Train Acc. (%)	Test Acc. (%)
0	99.125	98.5
1	98.725	99.3
2	96.925	96.7
3	96.8	95.9
4	97.7	97.1
5	95.925	94.5
6	98.15	97.6
7	98.025	96.9
8	95.25	94.6
9	95.475	96.1

Create 9 graphs (one for each label) with two ROC curves (one for each model).

project1 > notebook.ipynb

401:1 LF UTF-8 4 spaces ml_0135

Thu Apr 13 11:08 PM

project1 Version control

notebook.ipynb implementation.py

Add Configuration... Managed: http://localhost:8890 Python 3 (ipykernel) Trusted 10 28

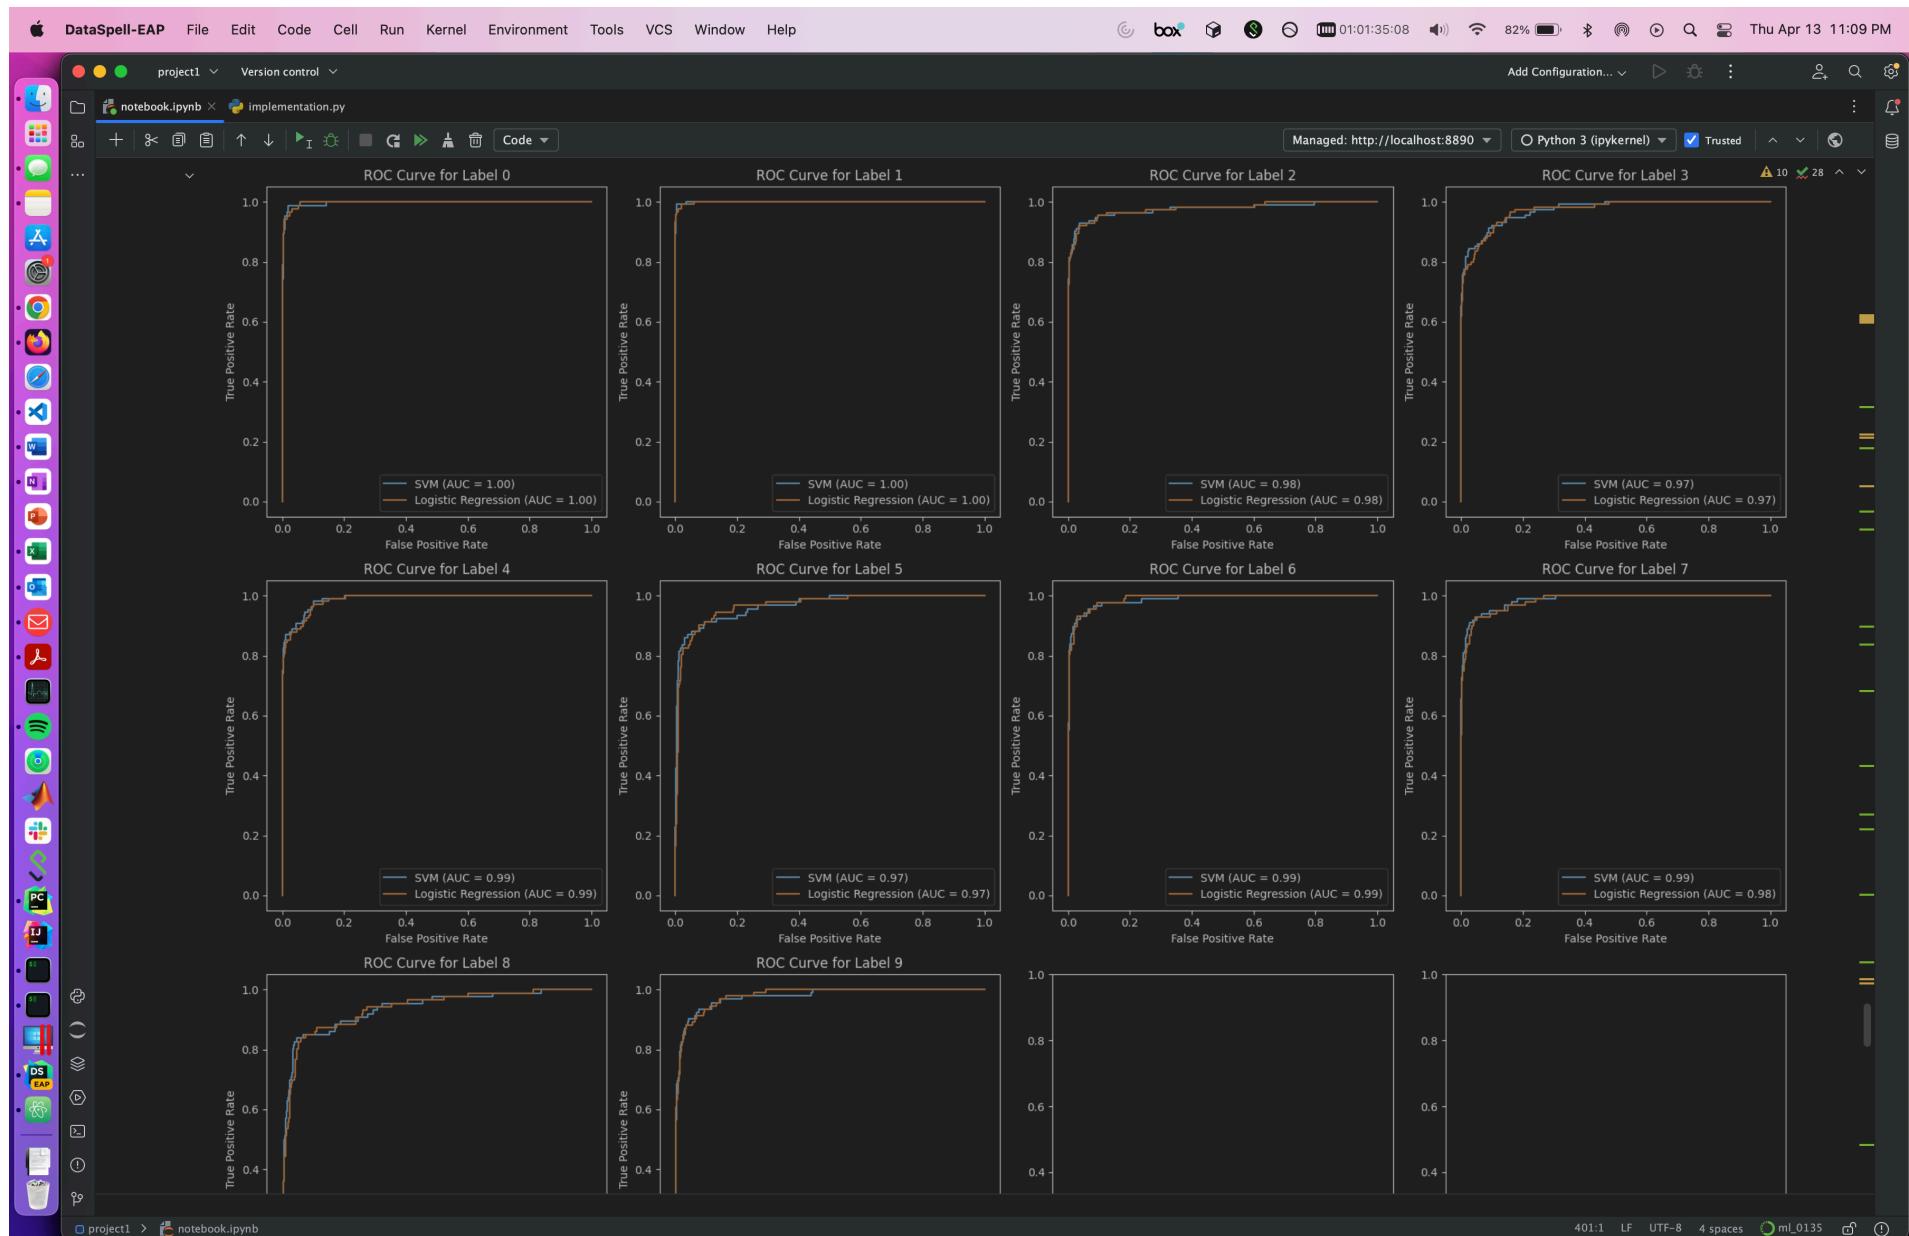
Create 9 graphs (one for each label) with two ROC curves (one for each model).

```
In 42 1 # TODO
2 from sklearn.metrics import roc_curve, auc
3
4 # Define the classes
5 classes = np.unique(train_samples_labels)
6
7 # Create subplots for each label
8 fig, axs = plt.subplots(3, 4, figsize=(20, 15))
9 axs = axs.ravel()
10 threshold = 0.0
11 # Iterate over each class
12 for i, curr_class in enumerate(classes):
13
14     # Extract the true labels for the current class
15     y_true = np.where(test_samples_labels == curr_class, 1, 0)
16
17     # Calculate the ROC curve and AUC for SVM model
18     fpr_svm, tpr_svm, threshold = roc_curve(y_true, test_probs_svm[i])
19     auc_svm = auc(fpr_svm, tpr_svm)
20
21     # Calculate the ROC curve and AUC for logistic regression model
22     fpr_lr, tpr_lr, threshold = roc_curve(y_true, test_probs_lr[i])
23     auc_lr = auc(fpr_lr, tpr_lr)
24
25     # Plot the ROC curves
26     axs[i].plot(fpr_svm, tpr_svm, label='SVM (AUC = {:.2f})'.format(auc_svm))
27     axs[i].plot(fpr_lr, tpr_lr, label='Logistic Regression (AUC = {:.2f})'.format(auc_lr))
28
29     # Set the axis labels and title
30     axs[i].set_xlabel('False Positive Rate')
31     axs[i].set_ylabel('True Positive Rate')
32     axs[i].set_title('ROC Curve for Label {}'.format(curr_class))
33
34     # Add the legend
35     axs[i].legend()
36
37 plt.tight_layout()
38 plt.show()
```

Executed in 1s, 13 Apr at 22:50:54

ROC Curve for Label 0 ROC Curve for Label 1 ROC Curve for Label 2 ROC Curve for Label 3

project1 > notebook.ipynb 401:1 LF UTF-8 4 spaces ml_0135



project1 Version control

notebook.ipynb implementation.py

Code

Managed: http://localhost:8890 Python 3 (ipykernel) Trusted

ROC Curve for Label 8

ROC Curve for Label 9

```

# BONUS (+5 points): Non-linear kernel
## Intuition Behind Kernels
The SVM classifier obtained by solving the convex Lagrange dual of the primal max-margin SVM formulation is as follows:

$$ f(x) = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i^\top \mathbf{x} + b $$

where $N$ is the number of support vectors.

If you know the intuition behind a linear discriminant function, the non-parametric SVM classifier above is very easy to understand. Instead of imagining the original features of each data point, consider a transformation to a new feature space where the data point has $N$ features, one for each support vector. The value of the $i^{th}$ feature is equal to the value of the kernel between the $i^{th}$ support vector and the data point is classified. The original (possibly non-linear) SVM classifier is like any other linear discriminant in this space.

Note that after the transformation, the original features of the data point are irrelevant. Its dot products with support vectors (special data points chosen by the SVM optimization algorithm) represent it only. One of my professors used a loose analogy while explaining this idea: A person has seen lakes, rivers, streams, fords, etc., but has never seen the sea. How would you explain to this person what a sea is? By relating the amount of water in an ocean to that found in a water body, the person already knows, etc.

In some instances, like the RBF kernel, defining the transformed features in terms of the original features of a data point leads to an infinite-dimensional representation. Unfortunately, though this is an awe-inspiring fact often mentioned while explaining how powerful SVMs are, it drops in only after repeated encounters with the idea ranging from introductory machine learning to statistical learning theory.

## Intuition Behind Gaussian Kernels

The Gaussian/RBF kernel is as follows:

$$ K(x, y) = \exp(-\frac{\|x-y\|^2}{2\sigma^2}) $$

Like any other kernel, we can understand the RBF kernel regarding feature transformation via the dot products given above. However, the intuition that helps best when analyzing the RBF kernel is that of the Gaussian

```

401:1 LF UTF-8 4 spaces ml_0135