



Kernel Methods for regression and classification



Many ideas/slides attributable to:
Dan Sheldon (U.Mass.)

James, Witten, Hastie, Tibshirani (ISL/ESL books)

Prof. Mike Hughes

Objectives for Day 19: Kernels

Big idea: Use kernel functions (similarity function with special properties) to obtain flexible high-dimensional feature transformations without explicit features

- From linear regression (LR) to kernelized LR
- What is a kernel function?
 - Basic properties
 - Example: Polynomial kernel
 - Example: Squared Exponential kernel
- Kernels for classification
 - Logistic Regression
 - SVMs


Task: Regression & Classification

Supervised
Learning

Unsupervised
Learning



Reinforcement
Learning

y is a numeric variable
e.g. sales in \$\$



Keys to Regression Success

- Feature transformation + linear model
- Penalized weights to avoid overfitting



(c) Alexander Ihler

Can fit **linear** functions to **nonlinear** features

A nonlinear function of x :

$$\hat{y}(x_i) = \theta_0 + \theta_1 x_i + \theta_2 x_i^2 + \theta_3 x_i^3$$

Can be written as a linear function of $\phi(x_i) = [1 \ x_i \ x_i^2 \ x_i^3]$

$$\hat{y}(x_i) = \sum_{g=1}^4 \theta_g \phi_g(x_i) = \theta^T \phi(x_i)$$

“Linear regression” means linear in the parameters (weights, biases)

Features can be arbitrary transforms of raw data

What feature transform to use?

- Anything that works for your data!

- sin / cos for periodic data

- polynomials for high-order dependencies

$$\phi(x_i) = [1 \ x_i \ x_i^2 \ \dots]$$

- interactions between feature dimensions

$$\phi(x_i) = [1 \ x_{i1}x_{i2} \ x_{i3}x_{i4} \ \dots]$$

- Many other choices possible

Review: Linear Regression

Prediction: Linear transform of G-dim features



$$\hat{y}(x_i, \theta) = \theta^T \phi(x_i) = \sum_{g=1}^G \theta_g \cdot \phi(x_i)_g$$

Training: Solve optimization problem

$$\min_{\theta} \sum_{n=1}^N (y_n - \hat{y}(x_n, \theta))^2 \quad + \text{L2 penalty (optional)}$$

Problems with high-dim features

- Feature transformation + linear model



(c) Alexander Ihler

How expensive is this transformation?
(Runtime and storage)

Thought Experiment

- Suppose that the optimal weight vector can be exactly constructed via a *linear combination* of the training set feature vectors

$$\theta^* = \alpha_1 \phi(x_1) + \alpha_2 \phi(x_2) + \dots + \alpha_N \phi(x_N)$$

Each alpha is a **scalar**

Each feature vector is a **vector of size G**

Justification?

Is optimal theta a linear combo of feature vectors?

Stochastic gradient descent, with 1 example per batch, can be seen as creating optimal weight vector of this form

- Starting with all zero vector
- In each step, adding a weight * feature vector

Each update step:

$$\theta_t \leftarrow \theta_{t-1} - \eta \cdot \frac{d}{d\theta} \text{loss}(y_n, \theta^T \phi(x_n))$$

*Let's simplify this via
chain rule!*

Justification?

Stochastic gradient descent, with 1 example per batch, can be seen as creating optimal weight vector of this form

- Starting with all zero vector
- In each step, adding a weight * feature vector

Each update step:

$$\theta_t \leftarrow \theta_{t-1} - \underset{\text{scalar}}{\eta} \cdot \underset{\text{scalar}}{\frac{d}{da} \text{loss}(y_n, a)} \cdot \underset{\text{Vector of size } G}{\frac{d}{d\theta} \theta^T \phi(x_n)}$$

Justification?

Stochastic gradient descent, with 1 example per batch, can be seen as creating optimal weight vector of this form

- Starting with all zero vector
- In each step, adding a weight * feature vector

Each update step:

$$\theta_t \leftarrow \theta_{t-1} - \underset{\text{scalar}}{\eta} \cdot \underset{\text{scalar}}{\frac{d}{da} \text{loss}(y_n, a)} \cdot \underset{\substack{\text{Vector of size } G \\ \text{(simplified)}}}{\phi(x_n)}$$

How to Predict in this thought experiment

$$\theta^* = \alpha_1 \phi(x_1) + \alpha_2 \phi(x_2) + \dots + \alpha_N \phi(x_N)$$

Prediction:

$$\hat{y}(x_i, \theta) = \theta^T \phi(x_i) :$$

$$\hat{y}(x_i, \theta^*) = \left(\sum_{n=1}^N \alpha_n \phi(x_n) \right)^T \phi(x_i)$$

How to Predict in this thought experiment

$$\theta^* = \alpha_1 \phi(x_1) + \alpha_2 \phi(x_2) + \dots + \alpha_N \phi(x_N)$$

Prediction:

$$\hat{y}(x_i, \theta) = \theta^T \phi(x_i) :$$

$$\hat{y}(x_i, \theta^*) = \sum_{n=1}^N \alpha_n \phi(x_n)^T \phi(x_i)$$

**Inner product
of test feature vector
with each training feature!**

Kernel Function

$$k(x_i, x_j) = \phi(x_i)^T \phi(x_j)$$

Input: any two vectors x_i and x_j

Output: scalar real

Interpretation: similarity function for x_i and x_j

Properties:

Larger output values mean i and j are more similar

Symmetric

Kernelized Linear Regression

- **Prediction:**

$$\hat{y}(x_i, \alpha, \{x_n\}_{n=1}^N) = \sum_{n=1}^N \alpha_n k(x_n, x_i) \\ = X$$

- **Training**

$$\min_{\alpha} \sum_{n=1}^N (y_n - \hat{y}(x_n, \alpha, X))^2$$

Can do all needed operations with only access to kernel (no feature vectors)

Compare: Linear Regression

Prediction: Linear transform of G-dim features

$$\hat{y}(x_i, \theta) = \theta^T \phi(x_i) = \sum_{g=1}^G \theta_g \cdot \phi(x_i)_g$$

Training: Solve optimization problem

$$\min_{\theta} \sum_{n=1}^N (y_n - \hat{y}(x_n, \theta))^2 \quad + \text{L2 penalty (optional)}$$

Why is kernel trick good idea?

Before:

Training problem optimized vector of size G

Prediction cost:

scales linearly with G (num. high-dim features)

After:

Training problem optimized vector of size N

Prediction cost:

scales linearly with N (num. train examples)
requires N evaluations of kernel

So we get some saving in runtime/storage if

G is bigger than N

AND we can compute k faster than inner product

Example: From Features to Kernels

$$x = [x_1 \quad x_2] \qquad z = [z_1 \quad z_2]$$

$$\phi(x) = [1 \quad x_1^2 \quad x_2^2 \quad \sqrt{2}x_1 \quad \sqrt{2}x_2 \quad \sqrt{2}x_1x_2]$$

$$k(x, z) = (1 + x_1z_1 + x_2z_2)^2$$

Compare:

What is relationship between these two functions defined above?

$$k(x, z) \qquad \phi(x)^T \phi(z)$$

Example: From Features to Kernels

$$x = [x_1 \quad x_2] \qquad z = [z_1 \quad z_2]$$

$$\phi(x) = [1 \quad x_1^2 \quad x_2^2 \quad \sqrt{2}x_1 \quad \sqrt{2}x_2 \quad \sqrt{2}x_1x_2]$$

$$k(x, z) = (1 + x_1z_1 + x_2z_2)^2$$

Compare:

What is relationship between these two functions defined above?

$$k(x, z) \quad = \quad \phi(x)^T \phi(z)$$

Punchline: Can sometimes find **faster** ways to compute high-dim. inner product

Cost comparison

$$x = [x_1 \quad x_2] \qquad z = [z_1 \quad z_2]$$

$$\phi(x) = [1 \quad x_1^2 \quad x_2^2 \quad \sqrt{2}x_1 \quad \sqrt{2}x_2 \quad \sqrt{2}x_1x_2]$$

$$k(x, z) = (1 + x_1z_1 + x_2z_2)^2$$

Compare:

Number of add and multiply ops to compute $\phi(x)^T \phi(z)$

Number of add and multiply ops to compute $k(x, z)$

Example:

Kernel cheaper than inner product

$$x = [x_1 \quad x_2]$$

$$\phi(x) = [1 \quad x_1^2 \quad x_2^2 \quad \sqrt{2}x_1 \quad \sqrt{2}x_2 \quad \sqrt{2}x_1x_2]$$

$$k(x, z) = (1 + x_1z_1 + x_2z_2)^2 \quad z = [z_1 \quad z_2]$$

Compare:

Number of add and multiply ops to compute $\phi(x)^T \phi(z)$

6 multiply and 5 add


Number of add and multiply ops to compute $k(x, z)$

3 multiply (include square) and 2 add

Squared Exponential Kernel

Assume x is a scalar

$$k(x, z) = e^{-(x-z)^2}$$



Also called “radial basis function (RBF)” kernel

Squared Exponential Kernel

Assume x is a scalar

$$\begin{aligned}k(x, z) &= e^{-(x-z)^2} \\&= e^{-x^2 - z^2 + 2xz} \\&= e^{-x^2} e^{-z^2} e^{2xz}\end{aligned}$$

Recall: Taylor series for e^x

$$e^x = \sum_{k=0}^{\infty} \frac{1}{k!} x^k = 1 + x + \frac{1}{2}x^2 + \dots$$

$$e^{2xz} = \sum_{k=0}^{\infty} \frac{2^k}{k!} x^k z^k$$


Squared Exponential Kernel

$$\begin{aligned}k(x, z) &= e^{-(x-z)^2} \\&= e^{-x^2 - z^2 + 2xz} \\&= e^{-x^2} e^{-z^2} \left(\sum_{k=0}^{\infty} \sqrt{\frac{2^k}{k!}} x^k \right) \left(\sum_{k=0}^{\infty} \sqrt{\frac{2^k}{k!}} z^k \right) \\&= \phi(x)^T \phi(z)\end{aligned}$$

Corresponds to an INFINITE DIMENSIONAL feature vector


$$\phi(x) = \left[\sqrt{\frac{2^0}{0!}} x^0 e^{-x^2} \quad \sqrt{\frac{2^1}{1!}} x^1 e^{-x^2} \quad \dots \quad \sqrt{\frac{2^k}{k!}} x^k e^{-x^2} \quad \dots \right]$$

Kernelized Regression Demo



Linear Regression

```
clf = sklearn.linear_model.LinearRegression()  
clf.fit(x_train, y_train)  
plot_model(x_test, clf)
```



Kernel Matrix for training set

- **$K : N \times N$ symmetric matrix**

$$K = \begin{bmatrix} k(x_1, x_1) & k(x_1, x_2) \dots k(x_1, x_N) \\ k(x_2, x_1) & k(x_2, x_2) \dots k(x_2, x_N) \\ \vdots & \\ k(x_N, x_1) & k(x_N, x_2) \dots k(x_N, x_N) \end{bmatrix}$$

Linear Regression with Kernel


100 training examples in `x_train`
505 test examples in `x_test`

```
def linear_kernel(X, Z):  
    '''  
    Compute dot product between each row of X and each row of Z  
    '''  
    m1, _ = X.shape  
    m2, _ = Z.shape  
    K = np.zeros((m1, m2))  
    for i in range(m1):  
        for j in range(m2):  
            K[i, j] = np.dot(X[i, :], Z[j, :])  
    return K  
  
K_train = linear_kernel(x_train, x_train) + 1e-10 * np.eye(N) # see note below  
K_test  = linear_kernel(x_test, x_train)  
  
print("Shape of K_train: %s" % str(K_train.shape))  
print("Shape of K_test: %s" % str(K_test.shape))
```


```
Shape of K_train: (100, 100)  
Shape of K_test: (505, 100)
```

Linear Regression with Kernel


```
clf = sklearn.linear_model.LinearRegression()  
clf.fit(K_train, y_train)  
plot_model(K_test, clf)
```




Polynomial Kernel, deg. 5



Polynomial Kernel, deg. 12



Gaussian kernel (aka sq. exp.)



Kernel Regression in sklearn

sklearn.kernel_ridge.KernelRidge

```
class sklearn.kernel_ridge. KernelRidge (alpha=1, kernel='linear', gamma=None, degree=3,  
coef0=1, kernel_params=None) \[source\]
```

```
fit (X, y=None, sample_weight=None)
```

[\[source\]](#)

Demo will use

kernel='precomputed'

Fit Kernel Ridge regression model

Parameters: **X :** *{array-like, sparse matrix}, shape = [n_samples, n_features]*

Training data. If kernel == "precomputed" this is instead a precomputed kernel matrix, shape = [n_samples, n_samples].

y : *array-like, shape = [n_samples] or [n_samples, n_targets]*

Target values

sample_weight : *float or array-like of shape [n_samples]*

Individual weights for each sample, ignored if None is passed.

Returns: **self :** *returns an instance of self.*

Can kernelize any linear model

Regression: Prediction

$$\hat{y}(x_i, \alpha, \{x_n\}_{n=1}^N) = \sum_{n=1}^N \alpha_n k(x_n, x_i)$$

Logistic Regression: Prediction

$$p(Y_i = 1|x_i) = \sigma(\hat{y}(x_i, \alpha, X))$$

Training for kernelized versions of

- * Linear Regression

- * Logistic Regression

$$\min_{\alpha} \sum_{n=1}^N (y_n - \hat{y}(x_n, \alpha, X))^2$$

$$\min_{\alpha} \sum_{n=1}^N \text{log_loss}(y_n, \sigma(\hat{y}(x_n, \alpha, X)))$$

SVMs: Prediction

$$\hat{y}(x_i) = w^T x_i + b$$

Make binary prediction via hard threshold

$$\begin{cases} 1 & \text{if } \hat{y}(x_i) \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

SVMs and Kernels: Prediction


$$\hat{y}(x_i) = \sum_{n=1}^N \alpha_n k(x_n, x_i)$$

Make binary prediction via hard threshold

$$\begin{cases} 1 & \text{if } \hat{y}(x_i) \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

Efficient training algorithms using modern quadratic programming
solve the dual optimization problem of SVM soft margin problem

Support vectors are often **small** fraction of all examples




Support vectors defined by **non-zero alpha** in kernel view


Data points i with non-zero weight α_i :

- Points with minimum margin (on optimized boundary)
- Points which violate margin constraint, but are still correctly classified
- Points which are misclassified

For all other training data, features have *no impact* on learned weight vector



SVM + Squared Exponential Kernel



Support vectors (green) for data separable by radial basis function kernels, and non-linear margin boundaries

Kernel Unit Objectives

Big idea: Use kernel functions (similarity function with special properties) to obtain flexible high-dimensional feature transformations without explicit features

- From linear regression (LR) to kernelized LR
- What is a kernel function?
 - Basic properties
 - Example: Polynomial kernel
 - Example: Squared Exponential kernel
- Kernels for classification
 - Logistic Regression
 - SVMs