


Binary Classification



Many slides attributable to:

Erik Sudderth (UCI)

Finale Doshi-Velez (Harvard)

James, Witten, Hastie, Tibshirani (ISL/ESL books)

Prof. Mike Hughes

Today's objectives (day 07)

Binary Classification Basics

- 3 steps of a classification task
 - Prediction
 - Predicting probabilities of each binary class
 - Making hard binary decisions
 - Training
 - Evaluation (much more in next class)
- A “taste” of 2 Methods
 - Logistic Regression
 - K-Nearest Neighbors

What will we learn?

Supervised
Learning

Unsupervised
Learning

Reinforcement
Learning

Training

Data, Label Pairs

$$\{x_n, y_n\}_{n=1}^N$$

Performance
measure

Task

data
 x

label
 y

Prediction

Evaluation

Before: Regression


Supervised
Learning

regression

Unsupervised
Learning

Reinforcement
Learning

y is a numeric variable
e.g. sales in \$\$



Task: Binary Classification


Supervised
Learning

**binary
classification**

Unsupervised
Learning

Reinforcement
Learning

y is a binary variable
(red or blue)



Example: Hotdog or Not



<https://www.theverge.com/tldr/2017/5/14/15639784/hbo-silicon-valley-not-hotdog-app-download>

Task: Multi-class Classification


Supervised
Learning

multi-class classification

Unsupervised
Learning

Reinforcement
Learning

y is a discrete variable
(red or blue or green or purple)



Binary Prediction Step

Goal: Predict label (0 or 1) given features x

- Input: $x_i \triangleq [x_{i1}, x_{i2}, \dots, x_{if} \dots x_{iF}]$
“features”
“covariates”
“predictors”
“attributes”
Entries can be real-valued, or other numeric types (e.g. integer, binary)
- Output: $y_i \in \{0, 1\}$ Binary label (0 or 1)
“responses”
“labels”

Binary Prediction Step

```
>>> # Given: pretrained binary classifier model
```

```
>>> # Given: 2D array of features x_NF
```

```
>>> x_NF.shape  
(N, F)
```

```
>>> yhat_N = model.predict(x_NF)
```

```
>>> yhat_N[:5] # peek at predictions  
[0, 0, 1, 0, 1]
```

```
>>> yhat_N.shape  
(N, )
```

Types of binary predictions

TN : true negative

FN : false negative

FP : false positive

TP : true positive

		classifier calls	
		"negative" C=0	"positive" C=1
true outcome	Y=0	TN	FP
	Y=1	FN	TP

Probability Prediction Step

Goal: Predict probability of event $y=1$ given features x

- Input: $x_i \triangleq [x_{i1}, x_{i2}, \dots, x_{if} \dots x_{iF}]$
“features”
“covariates”
“predictors”
“attributes”
Entries can be real-valued, or other numeric types (e.g. integer, binary)
- Output: \hat{p}_i
“probabilities”
Probability between 0 and 1
e.g. 0.001, 0.513, 0.987

Probability Prediction Step

```
>>> # Given: pretrained regression object model
```

```
>>> # Given: 2D array of features x_NF
```

```
>>> x_NF.shape  
(N, F)
```

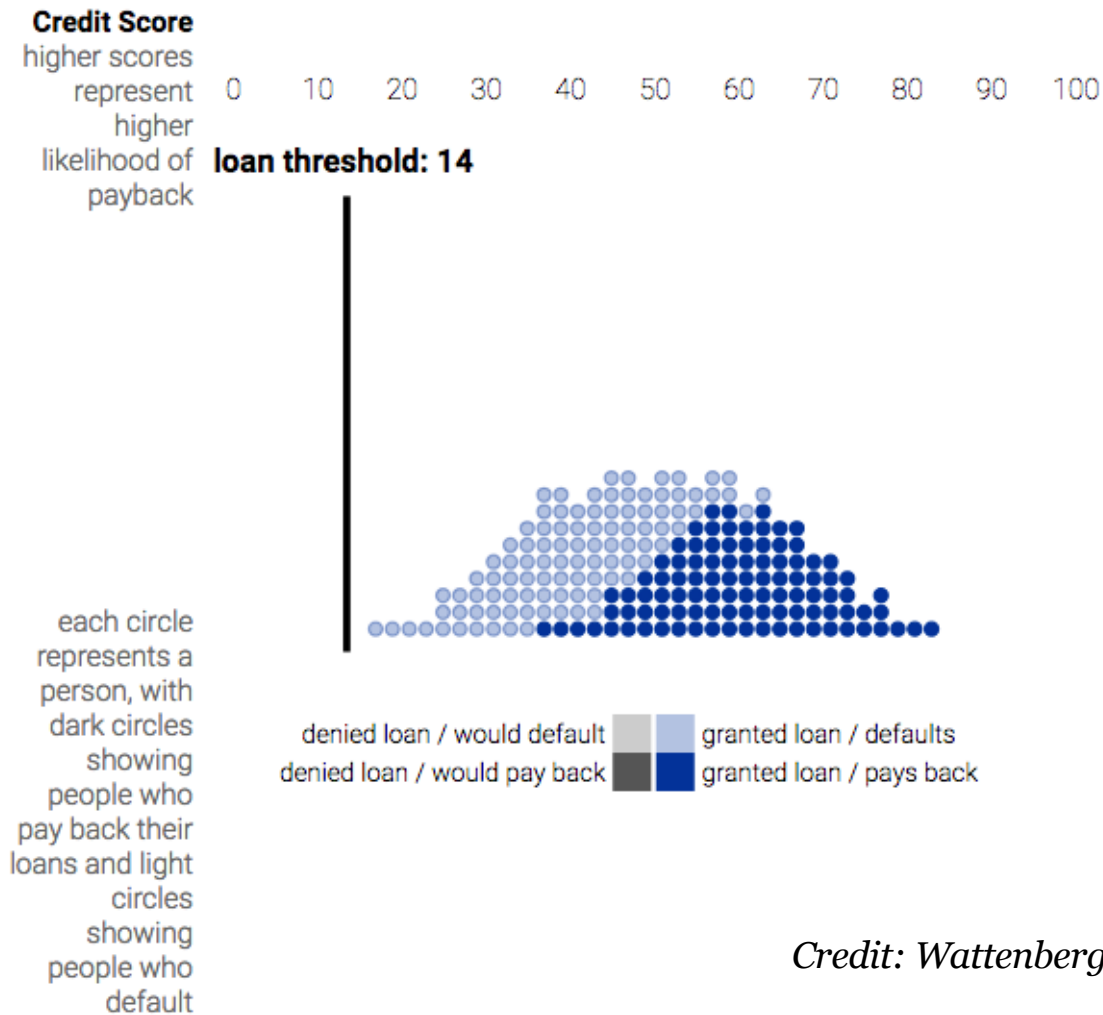
```
>>> yproba_N2 = model.predict_proba(x_NF)
```

```
>>> yproba_N2.shape  
(N, 2)
```

*Column index 1 gives
probability of positive label
given input features*

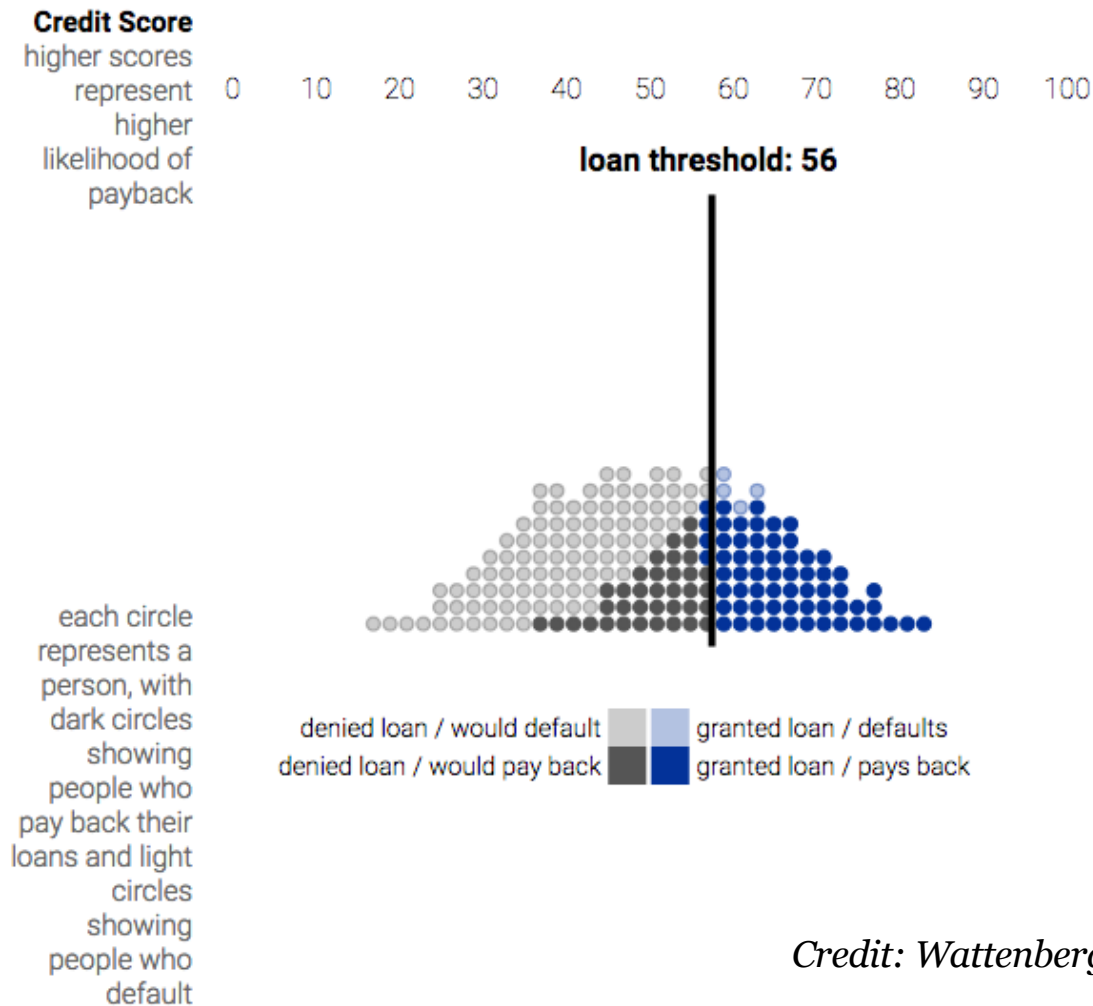
```
>>> yproba_N2[:, 1]  
[0.003, 0.358, 0.987, 0.111, 0.656]
```

Thresholding to get Binary Decisions



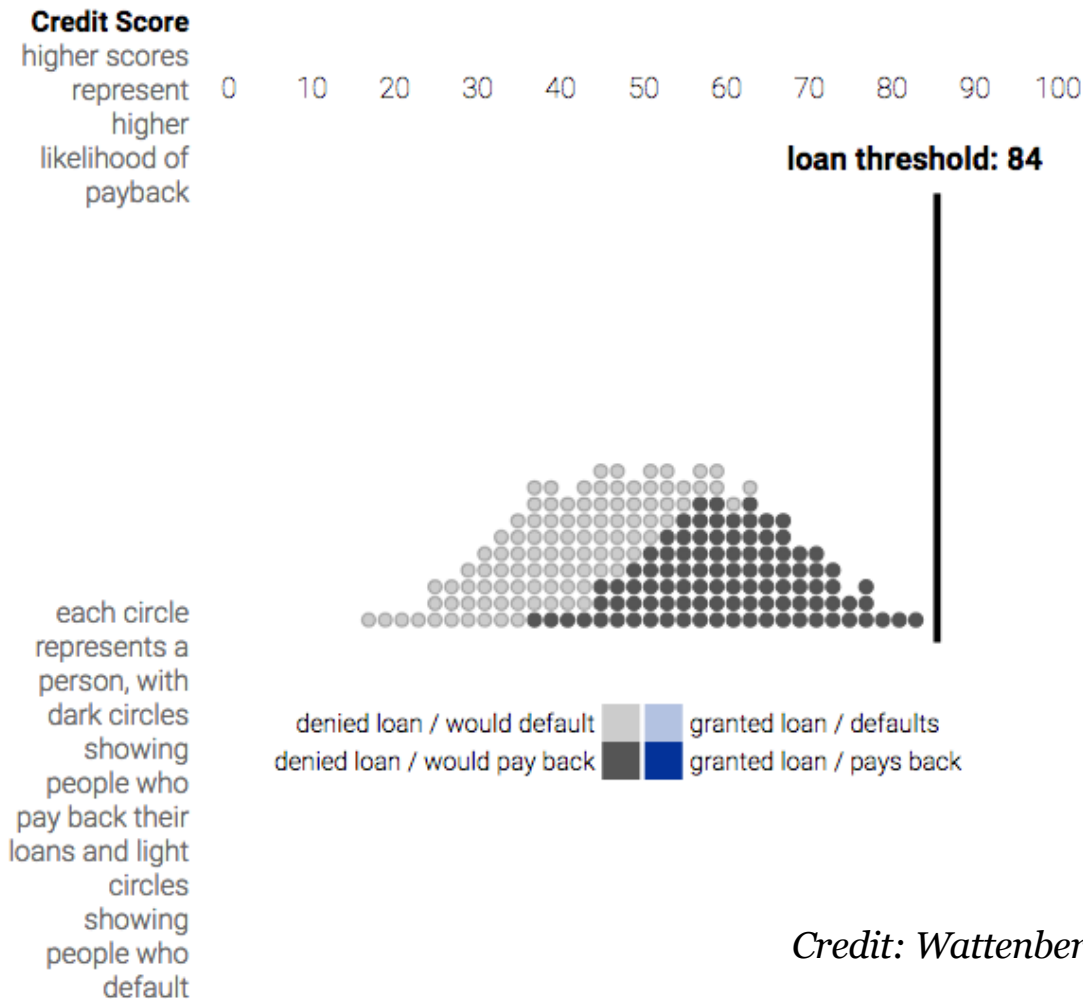
Credit: Wattenberg, Viégas, Hardt

Thresholding to get Binary Decisions



Credit: Wattenberg, Viégas, Hardt

Thresholding to get Binary Decisions



Credit: Wattenberg, Viégas, Hardt

Classifier: Training Step

Goal: Given a labeled dataset, learn a **function** that can perform (probabilistic) prediction well

- Input: Pairs of features and labels/responses

$$\{x_n, y_n\}_{n=1}^N$$

- Output: $\hat{y}(\cdot) : \mathbb{R}^F \rightarrow \{0, 1\}$

Useful to break into two steps:

- 1) Produce real-valued scores OR probabilities in $[0, 1]$*
- 2) Threshold to make binary decisions*

Classifier: Training Step

```
>>> # Given: 2D array of features x_NF
>>> # Given: 1D array of binary labels y_N

>>> y_N.shape
(N, )
>>> x_NF.shape
(N, F)

>>> model = BinaryClassifier()
>>> model.fit(x_NF, y_N)
>>> # Now can call predict or predict_proba
```

Classifier: Evaluation Step

Goal: Assess quality of predictions

Many ways in practice:

- 1) Evaluate probabilities / scores directly
cross entropy loss (aka log loss), hinge loss, ...
- 2) Evaluate binary decisions at specific threshold
accuracy, TPR, TNR, PPV, NPV, etc.
- 3) Evaluate across range of thresholds
ROC curve, Precision-Recall curve

Metric: Confusion Matrix

Counting **mistakes** in binary predictions

#TN : num. true negative

#FN : num. false negative

#TP : num. true positive

#FP : num. false positive

		classifier calls	
		"negative" C=0	"positive" C=1
true outcome	Y=0	#TN	#FP
	Y=1	#FN	#TP

Metric: Accuracy

accuracy = fraction of correct predictions

$$= \frac{TP + TN}{TP + TN + FN + FP}$$

Potential problem:

Suppose your dataset has 1 positive example and 99 negative examples

What is the accuracy of the classifier that always predicts "negative"?

Metric: Accuracy

accuracy = fraction of correct predictions

$$= \frac{TP + TN}{TP + TN + FN + FP}$$

Potential problem:

Suppose your dataset has 1 positive example and 99 negative examples

What is the accuracy of the classifier that always predicts "negative"?


99%!

Objectives: Classifier Overview

- 3 steps of a classification task
 - Prediction
 - Making hard binary decisions
 - Predicting class probabilities
 - Training
 - Evaluation
- A “taste” of 2 Methods
 - Logistic Regression
 - K-Nearest Neighbors

Logistic Sigmoid Function

Goal: Transform real values into probabilities



Logistic Regression

Parameters:

weight vector $w = [w_1, w_2, \dots w_f \dots w_F]$

bias scalar b

Prediction:

$$\hat{p}(x_i, w, b) = p(y_i = 1|x_i) \triangleq \text{sigmoid} \left(\sum_{f=1}^F w_f x_{if} + b \right)$$

Training:


find weights and bias that minimize “loss”

Measuring prediction quality for a probabilistic classifier

Use the log loss (aka “binary cross entropy”)

```
from sklearn.metrics import log_loss
```

$$\text{log_loss}(y, \hat{p}) = -y \log \hat{p} - (1 - y) \log(1 - \hat{p})$$



Advantages:

- smooth
- easy to take derivatives!

Logistic Regression: Training

Optimization: Minimize total log loss on train set


$$\min_{w,b} \sum_{n=1}^N \text{log_loss}(y_n, \hat{p}(x_n, w, b))$$

Algorithm: Gradient descent

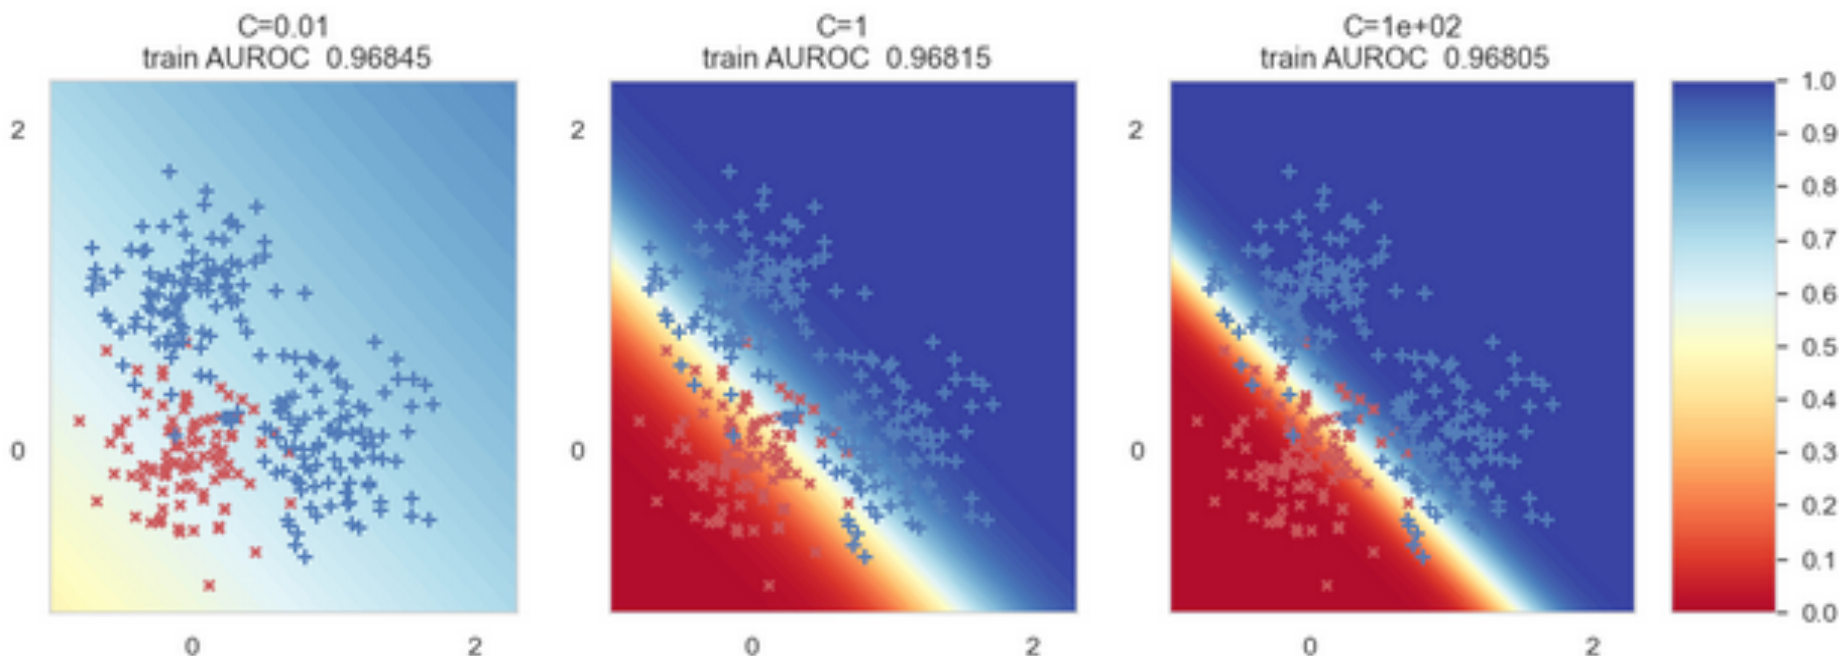

Avoid overfitting: Use L2 or L1 penalty on weights

Much more in depth in next class

Visualizing predicted probas for Logistic Regression



Visualizing predicted probas for Logistic Regression



Nearest Neighbor Classifier

Parameters:

none

Prediction:

- find “nearest” training vector to given input x
- predict y value of this neighbor

Training:

none needed (use training data as lookup table)

K nearest neighbor classifier

Parameters:

K : *number of neighbors*


Prediction:

- find K “nearest” training vectors to input x
- predict: vote most common y in neighborhood
- predict_proba: report fraction of labels

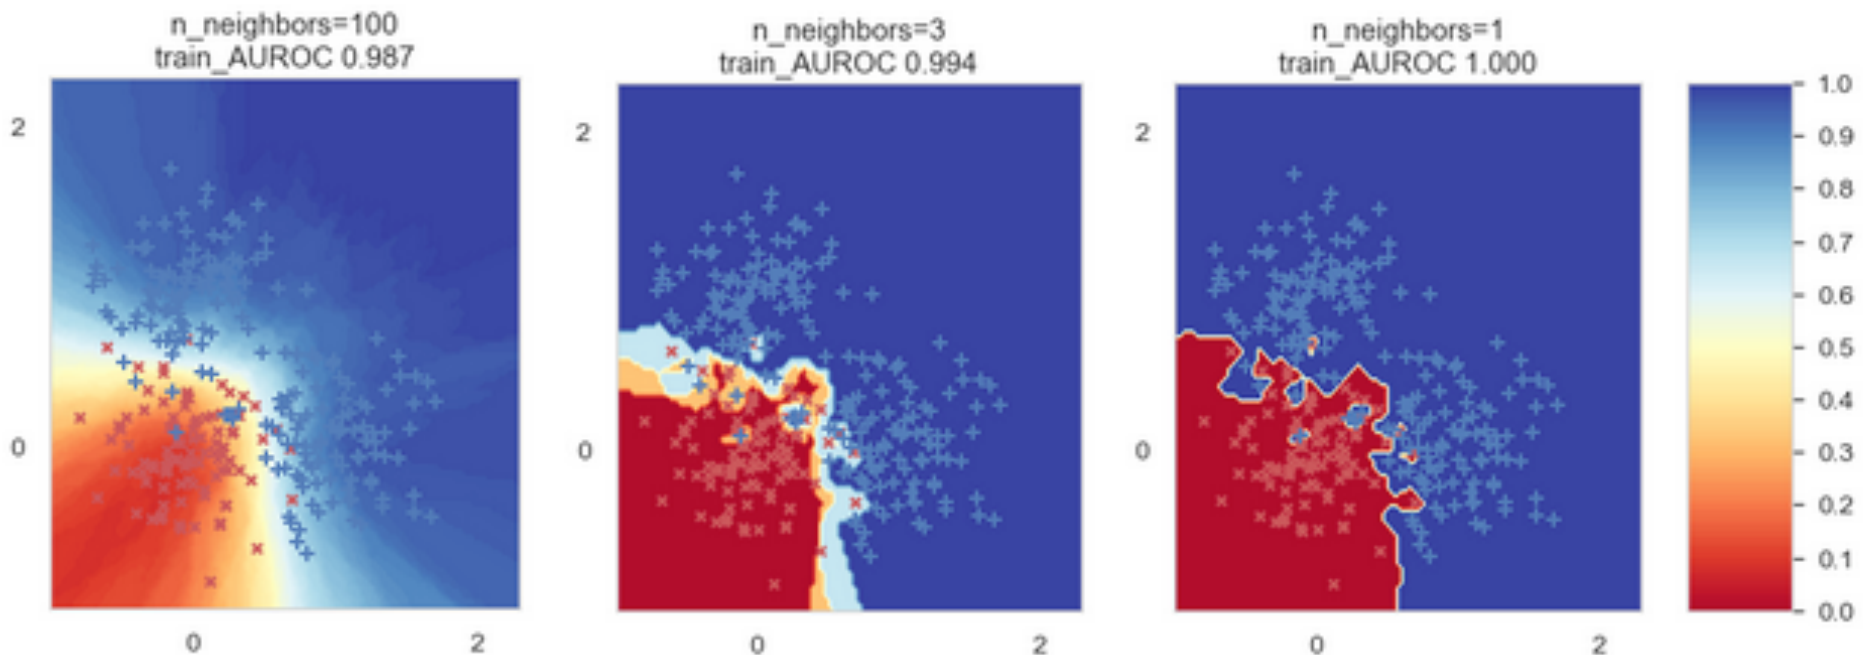

Training:

none needed (use training data as lookup table)

Visualizing predicted probas for K-Nearest Neighbors



Visualizing predicted probas for K-Nearest Neighbors



Summary of Methods

	Function class flexibility	Hyperparameters to select (control complexity)	Interpret?
Logistic Regression	Linear	L2/L1 penalty on weights	Inspect weights
K Nearest Neighbors Classifier	Piecewise constant	Number of Neighbors Distance metric	Inspect neighbors