A Deep Understanding of CNN

Joseph Robinson

Advisor: Yun (Raymond) Fu

ECE7399
Spring 2018

So What-Who Cares???

- Pre-Deep learning: saturation on existing datasets
  - Needed larger, more complex datasets to properly represent real-world data, in part, triggered Big Data Era
  - Need systems for Big Data, i.e., sample count & complexity
- Feature extraction, representation, and modeling were hand-crafted by experts, i.e., better if end-to-end learning
- Today's HW was not used to full capacity (i.e., GPGPU)
- Data-driven modern day demands *off-the-shelf* capability for a broader audience (i.e., accessible to those without PhD)
- Advances in 2012 brought deep learning center stage
- Toolboxes for deep learning provide variety of interfaces, level of implementation, and capability

Towards a Deeper Understanding of NN

Problem Definition

- Capture full complexity of big data without overfitting
- Maximize full capacity of modern day HW
- End-to-end training, opposed to professional with expert knowledge for data, hand-crafted features, and modeling

CNNs: Main Components

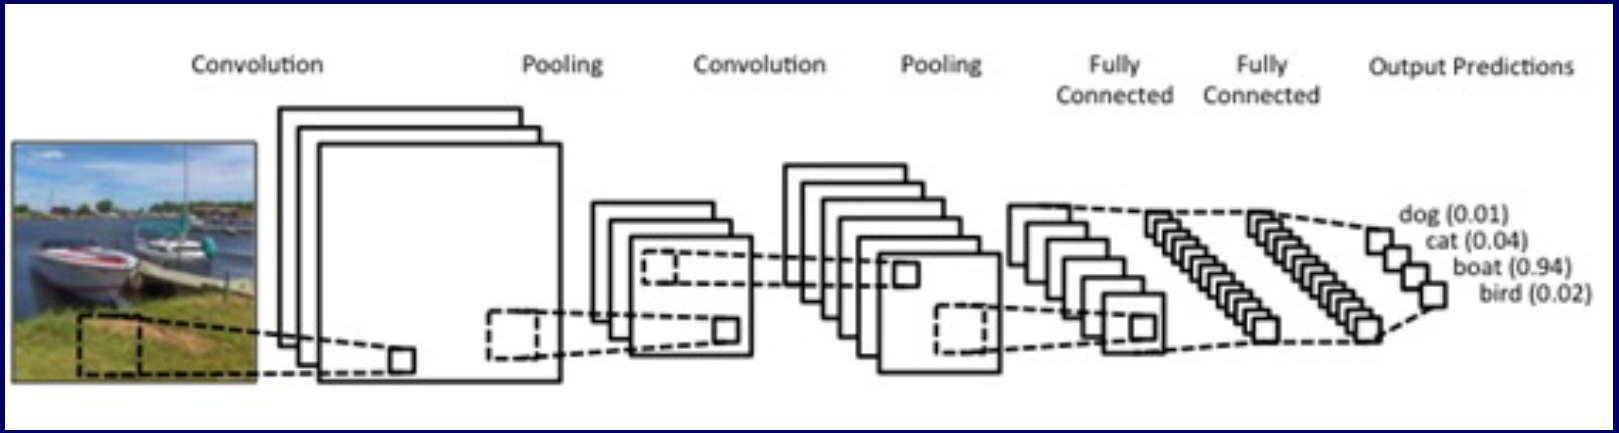

Neural Networks are constructed using variants of the same underlying components

1. Convolution Layer (Conv-Layer)
2. Non-Linear Activation
3. Pooling (i.e., sub-sampling)
4. Classification-Layer

CNNs: Main Components

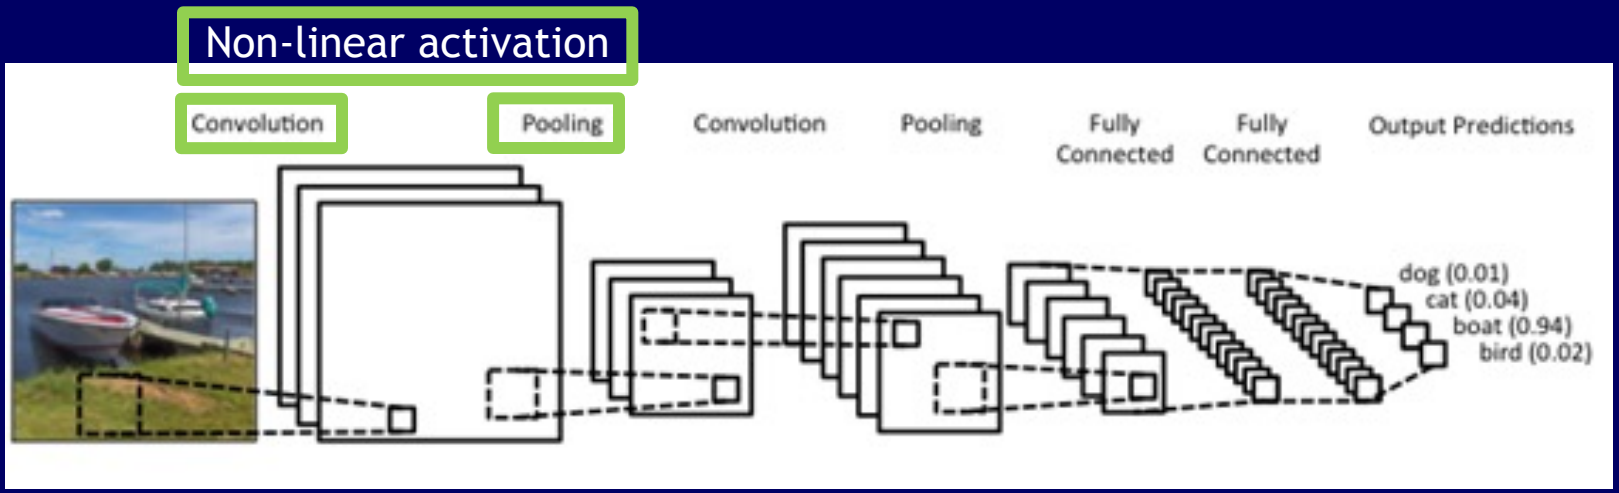

Neural Networks are constructed using variants of the same underlying components

1. Convolution Layer (Conv-Layer)
2. Non-Linear Activation
3. Pooling (i.e., sub-sampling)
4. Classification-Layer

CNNs: Main Components

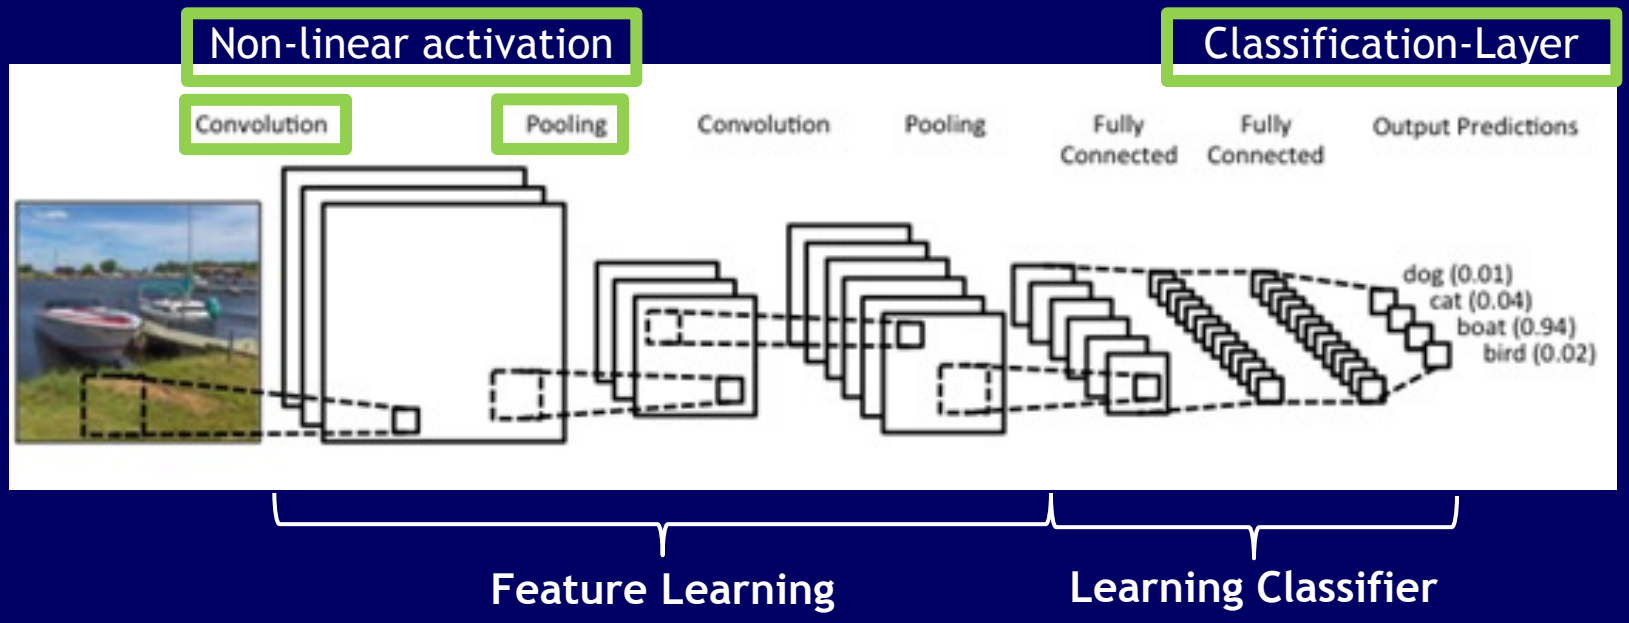

Neural Networks are constructed using variants of the same underlying components

1. Convolution Layer (Conv-Layer)
2. Non-Linear Activation
3. Pooling (i.e., sub-sampling)
4. Classification-Layer

Conv-Net: Conv-Layers

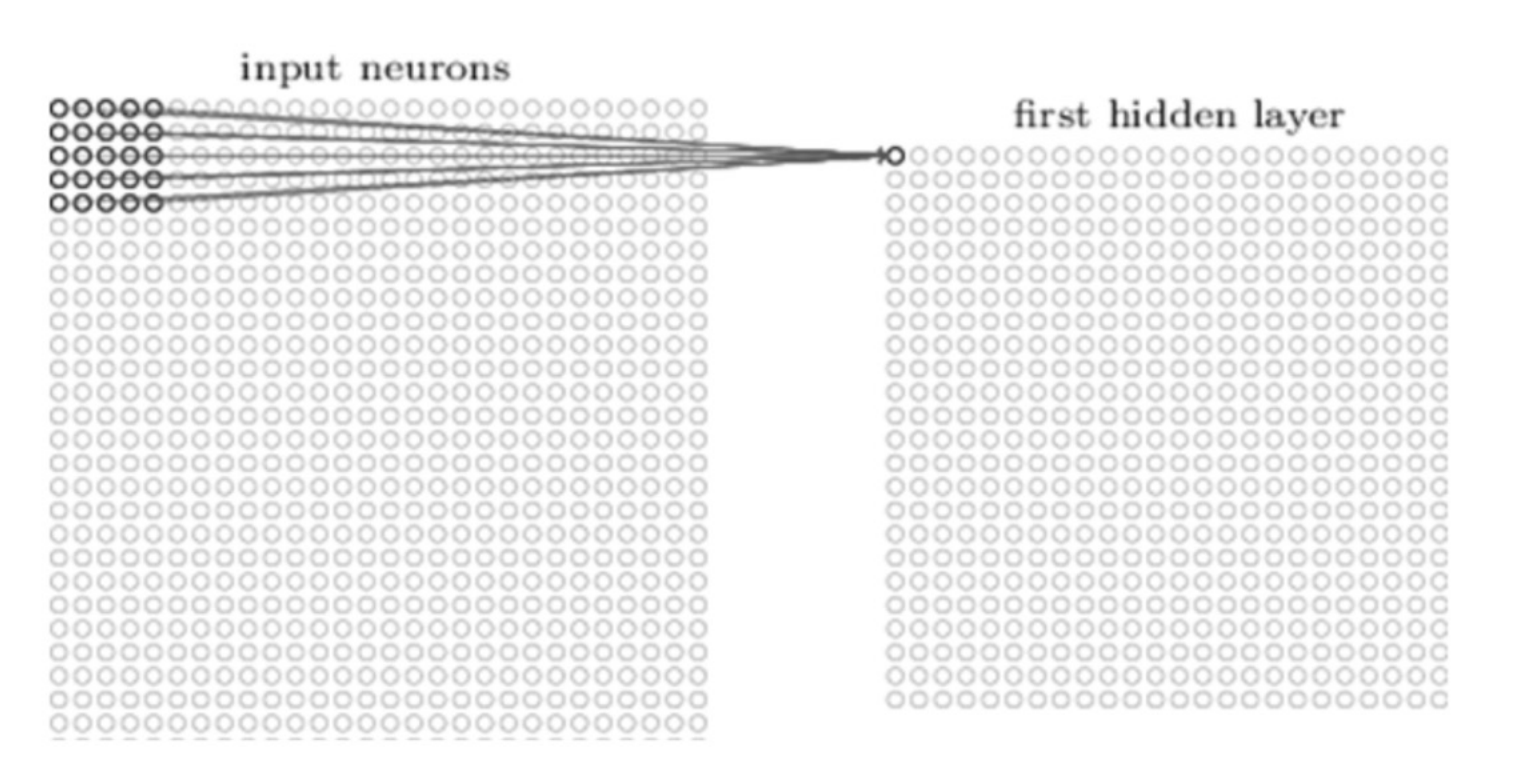

- Element-wise multiplication applied to identical copies of each neuron (aka kernel aka filter)
- Each neurons learn weight for different spatial locations
- Essentially, just a measure of correlation

Conv-Net: Conv-Layers

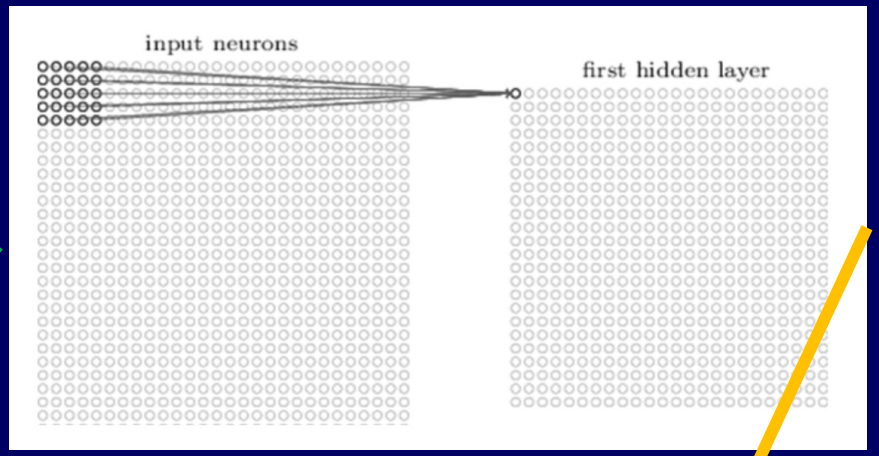

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Signal
(image or activations of previous layer)

1	0	1
0	1	0
1	0	1

Filter (i.e., units in hidden layer)

Convolution Layer (Continued)

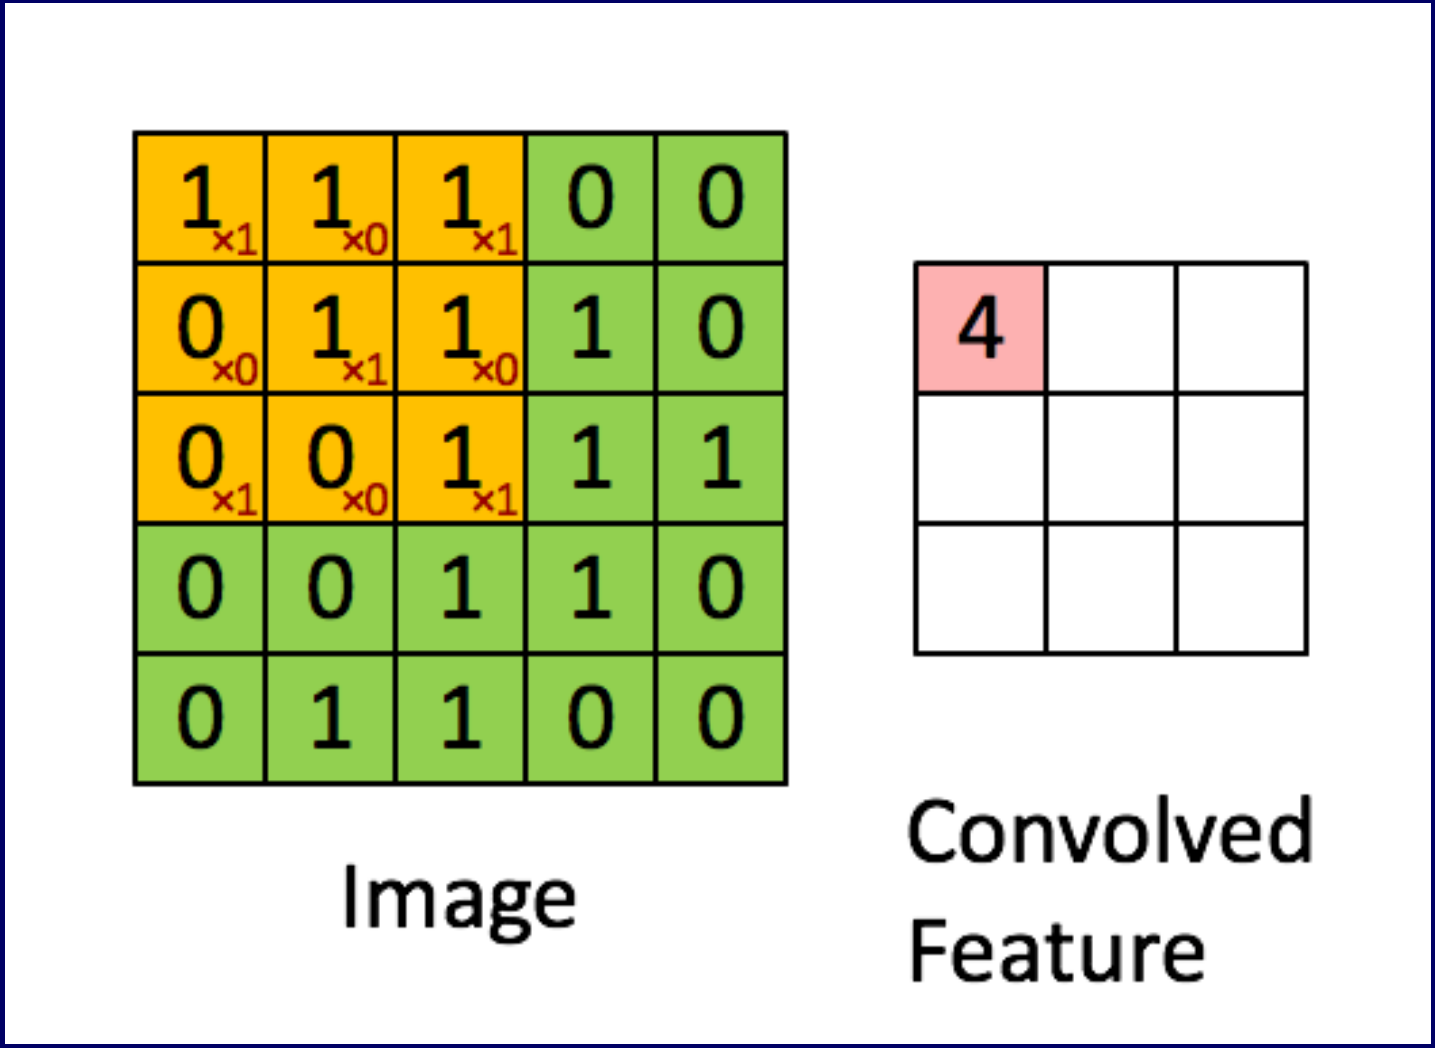

Convolution operation– Output matrix is called Convolved Feature or Feature Map

Convolution Layer (Continued)

Conv-Net: Conv-Layers

Convolution operation– Output matrix is called Convolved Feature or Feature Map

Convolution Layer (Continued)

Conv-Net: Conv-Layers

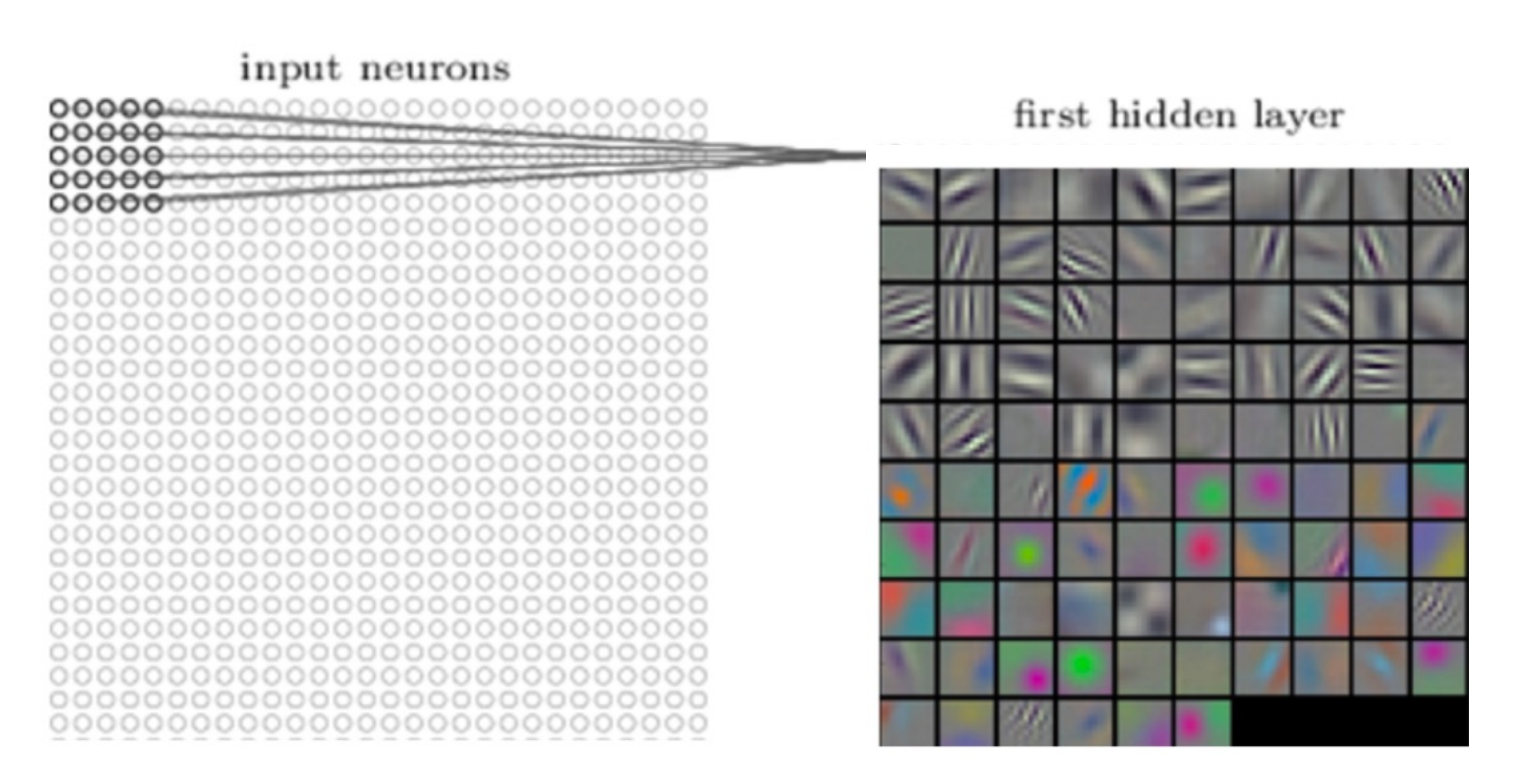

- Element-wise multiplication applied to identical copies of each neuron (aka kernel aka filter)
- Each neurons learn weight for different spatial locations
- Essentially, just a measure of correlation

Conv-Net: Activation Function

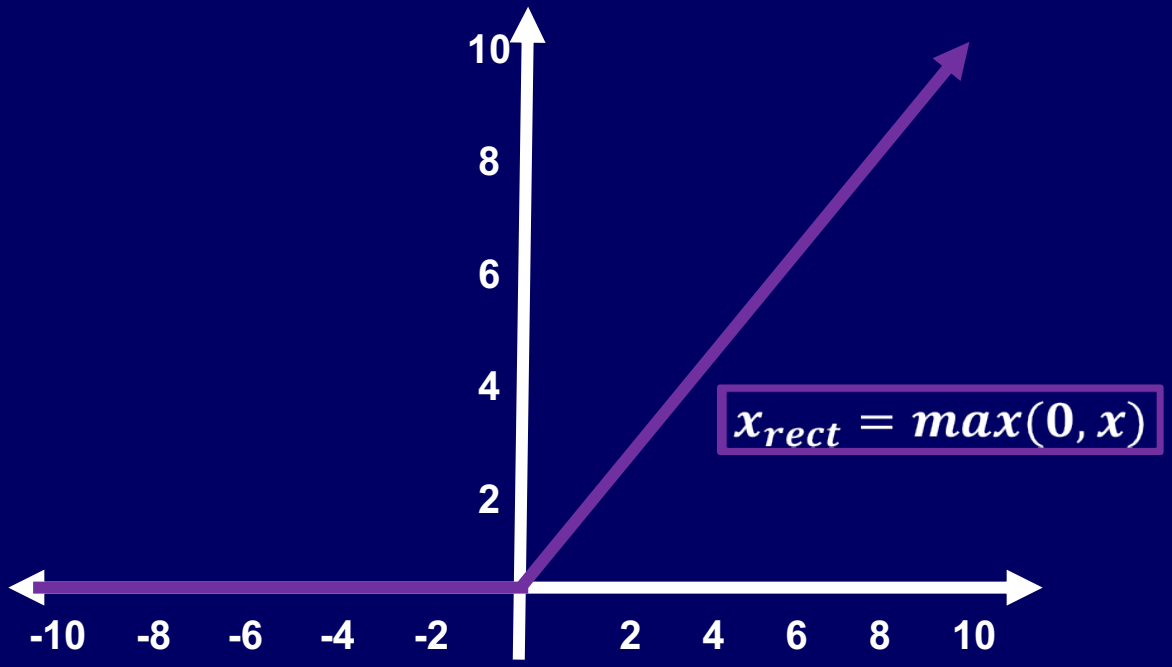

- Introduces non-linearity to representation (i.e., else, conv-layer is linear and unable to capture complex features)
- Just a threshold (e.g., ReLU sets negative values to zero)

Conv-Net: Activation Function

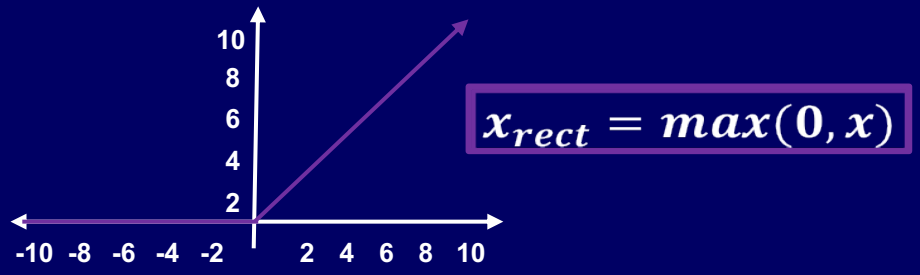

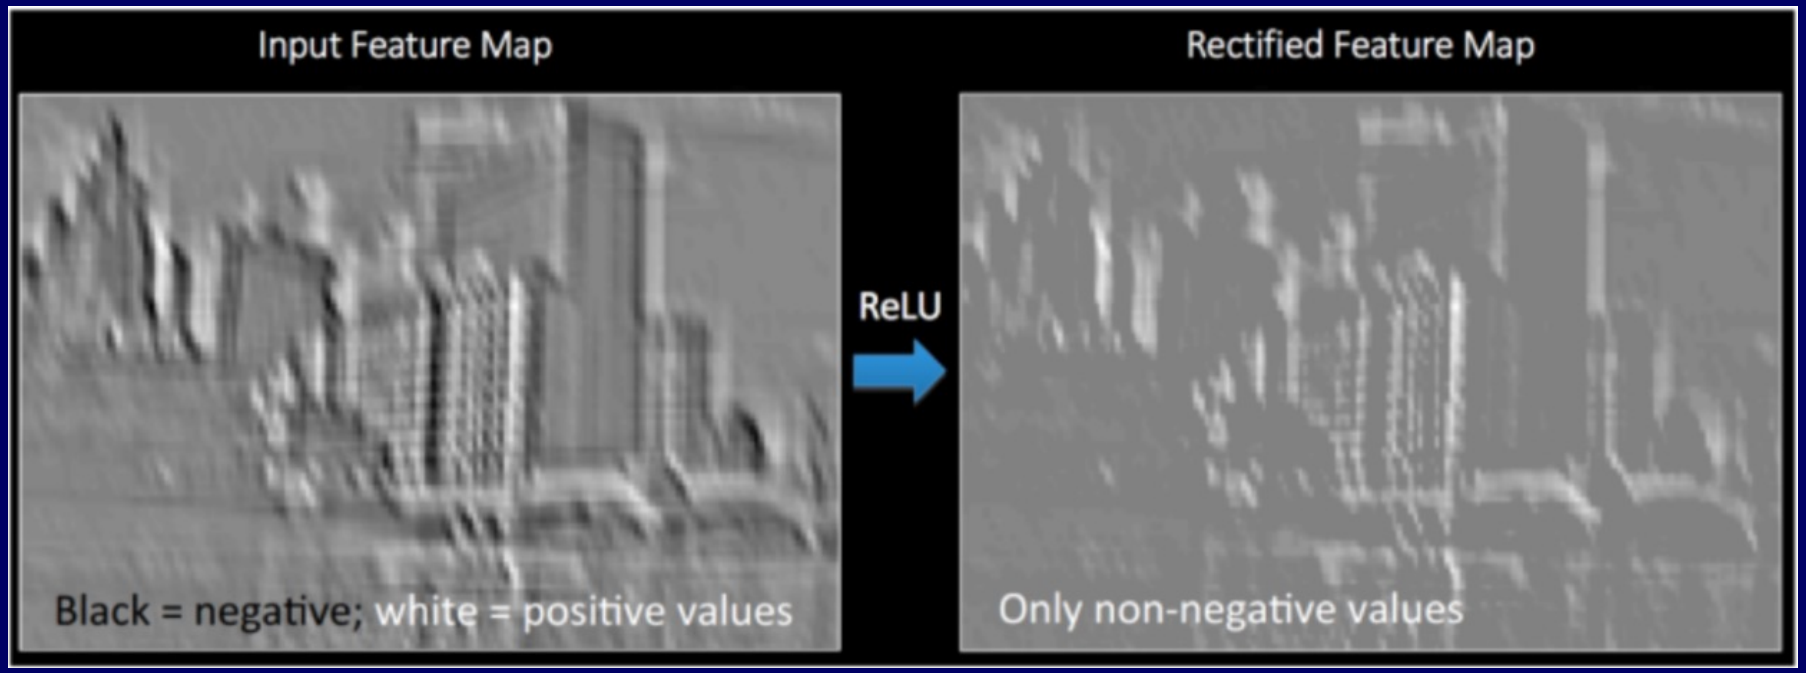

- Introduces non-linearity to representation (i.e., else, conv-layer is linear and unable to capture complex features)
- Just a threshold (e.g., ReLU sets negative values to zero)

Conv-Net: Subsampling (pooling or down-sampling)

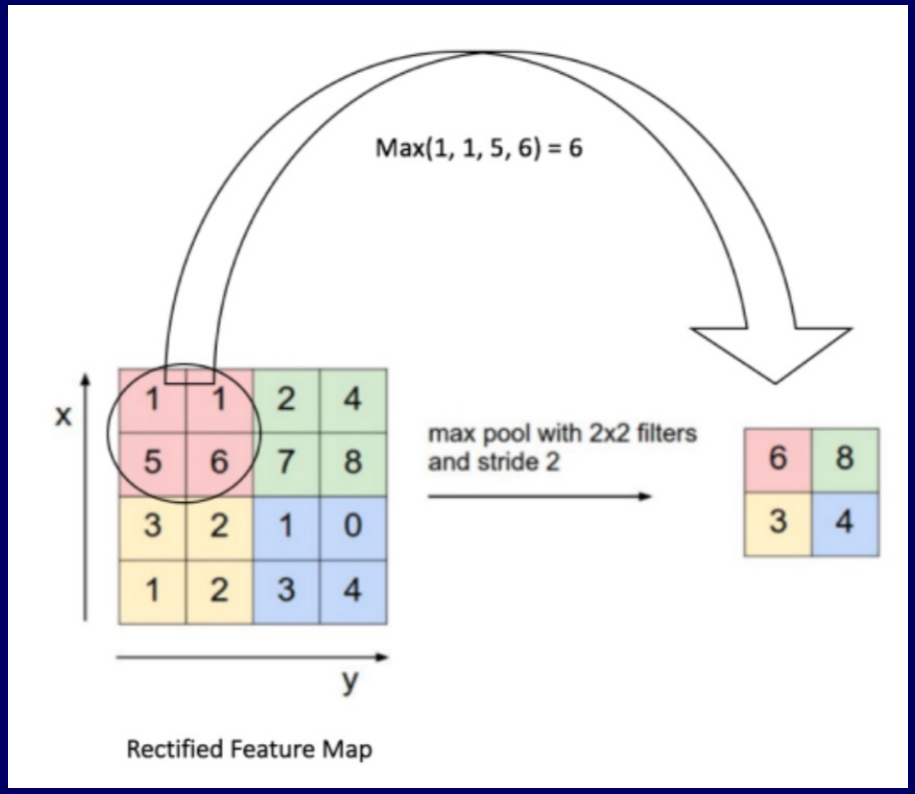

- Reduces dimension of feature map
- Introduces shift invariance

CNNs: Hyper-parameters

- Hyper-parameters: Parameters of parameters, manually set (i.e., typically optimized via cross validation, trial and error, or, less often, transfer expert knowledge from another visual task solved with deep learning)

CNNs: Hyper-parameters

- Hyper-parameters: Parameters of parameters, manually set (i.e., typically optimized via cross validation, trial and error, or, less often, transfer expert knowledge from another visual task solved with deep learning)
  1. **Depth**. Number of filters used at a given layer.

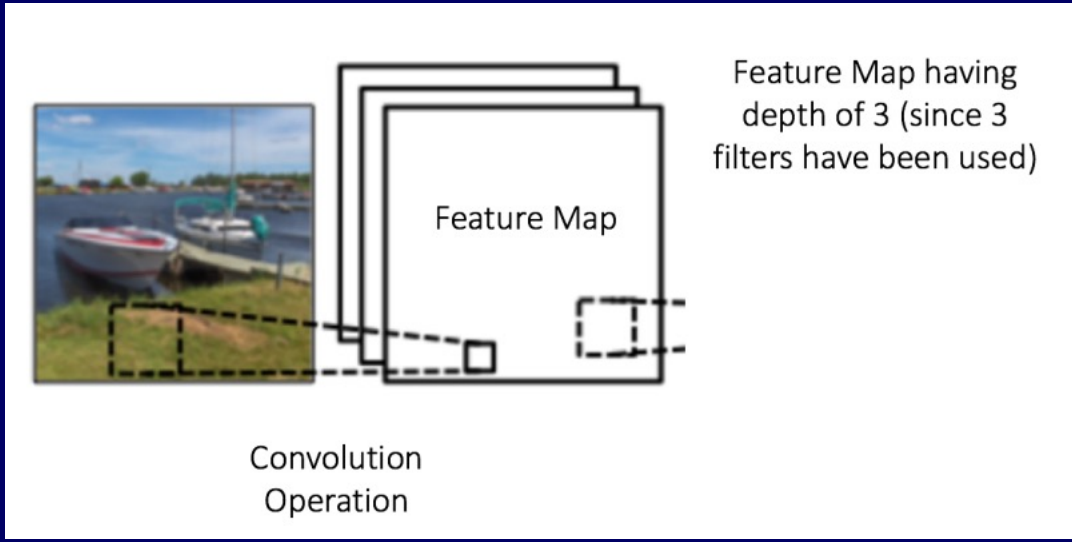

CNNs: Hyper-parameters

- Hyper-parameters: Parameters of parameters, manually set (i.e., typically optimized via cross validation, trial and error, or, less often, transfer expert knowledge from another visual task solved with deep learning)
  1. **Depth**. Number of filters used at a given layer.
  2. **Stride**. Number of neurons to skip between filters. Double the stride results in ½ sized feature maps.

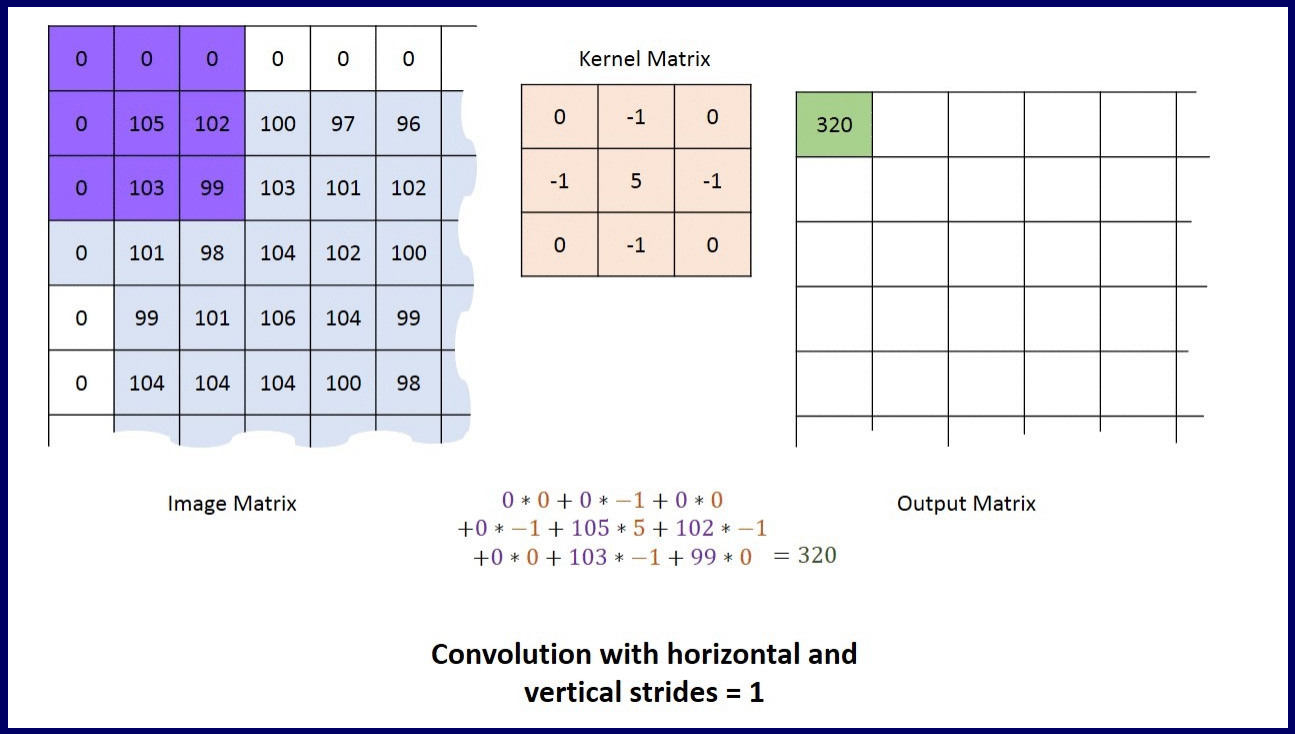

CNNs: Hyper-parameters

- Hyper-parameters: Parameters of parameters, manually set (i.e., typically optimized via cross validation, trial and error, or, less often, transfer expert knowledge from another visual task solved with deep learning)
  1. **Depth**. Number of filters used at a given layer.
  2. **Stride**. Number of neurons to skip between filters. Double the stride results in ½ sized feature maps.
  3. **Zero padding.** Padding borders with zeros is a means of tailoring edge values. Moreover, it provides a way of controlling size of output.

Zero-Padded

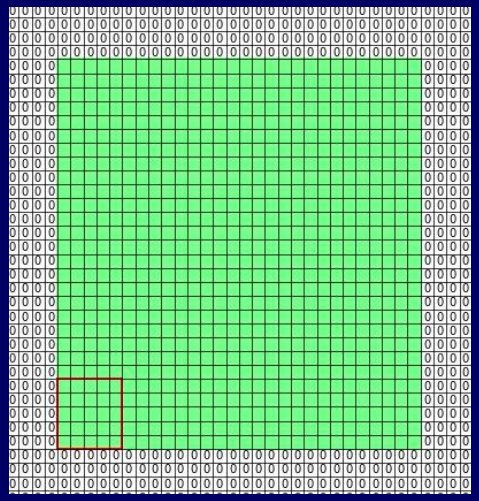

No Zero-Pad

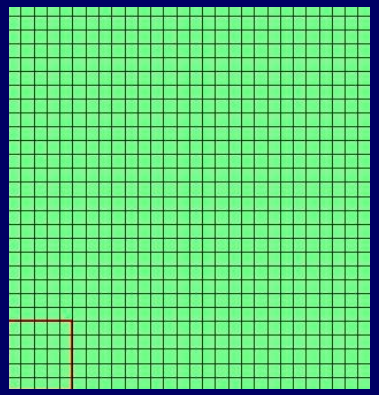

CNNs: Feed Forward

Feed-Forward

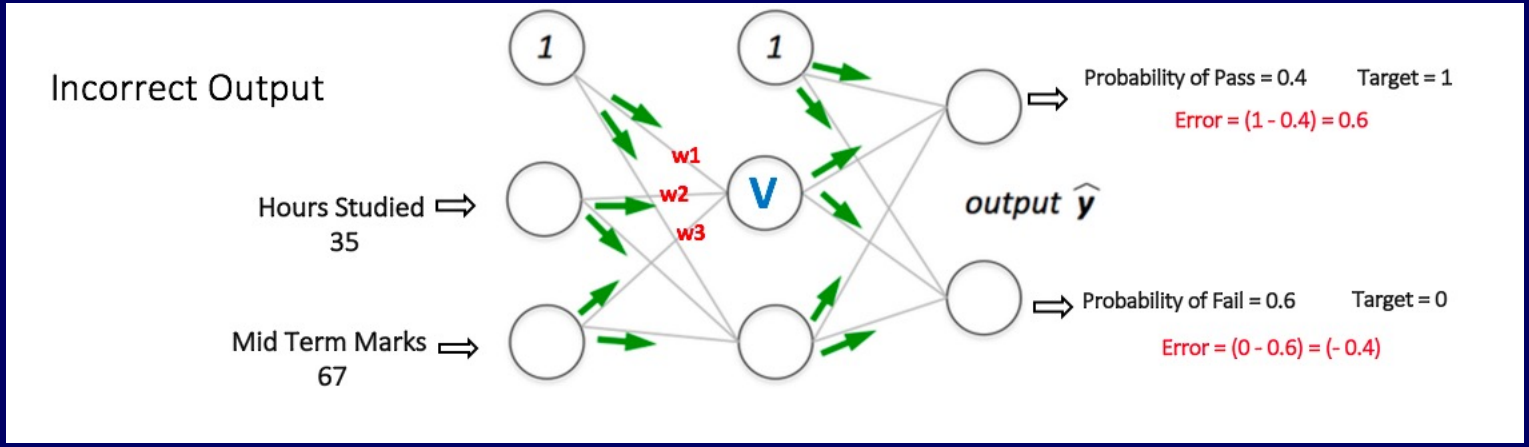

Back Propagation

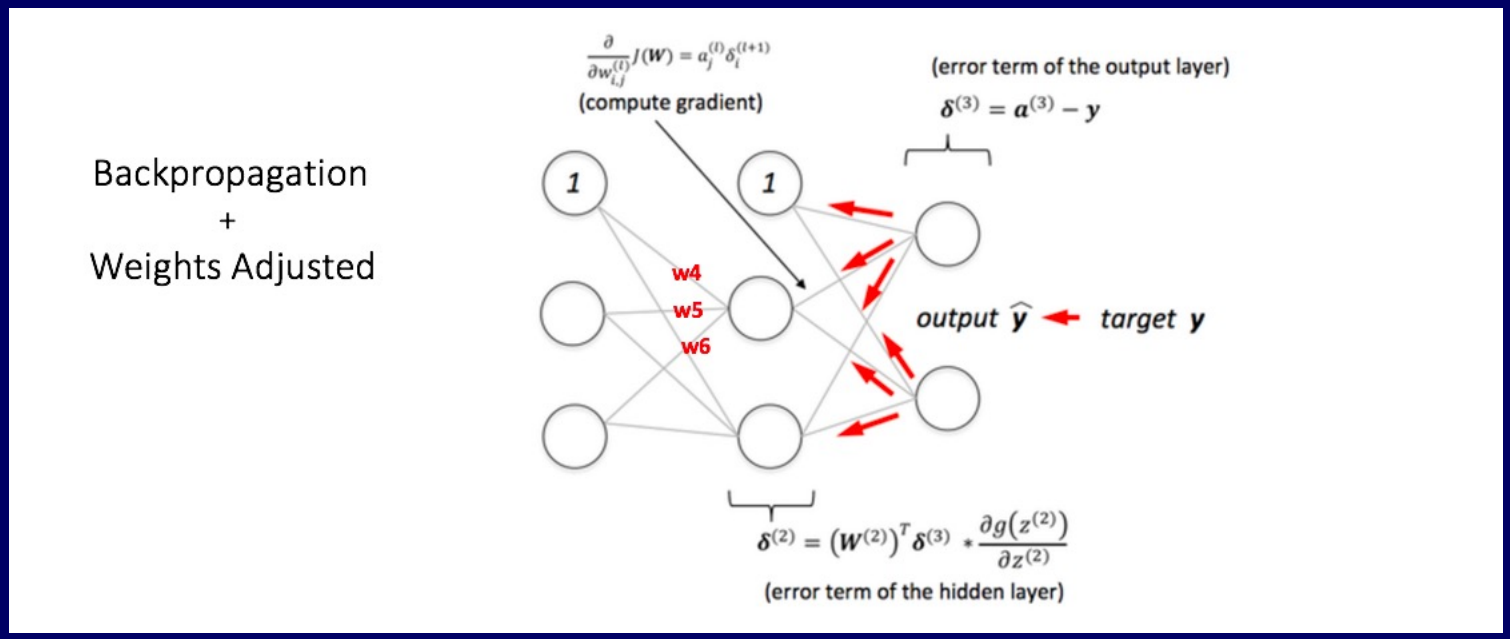

What is learned at Different Levels

Conv-Net: Learned Features

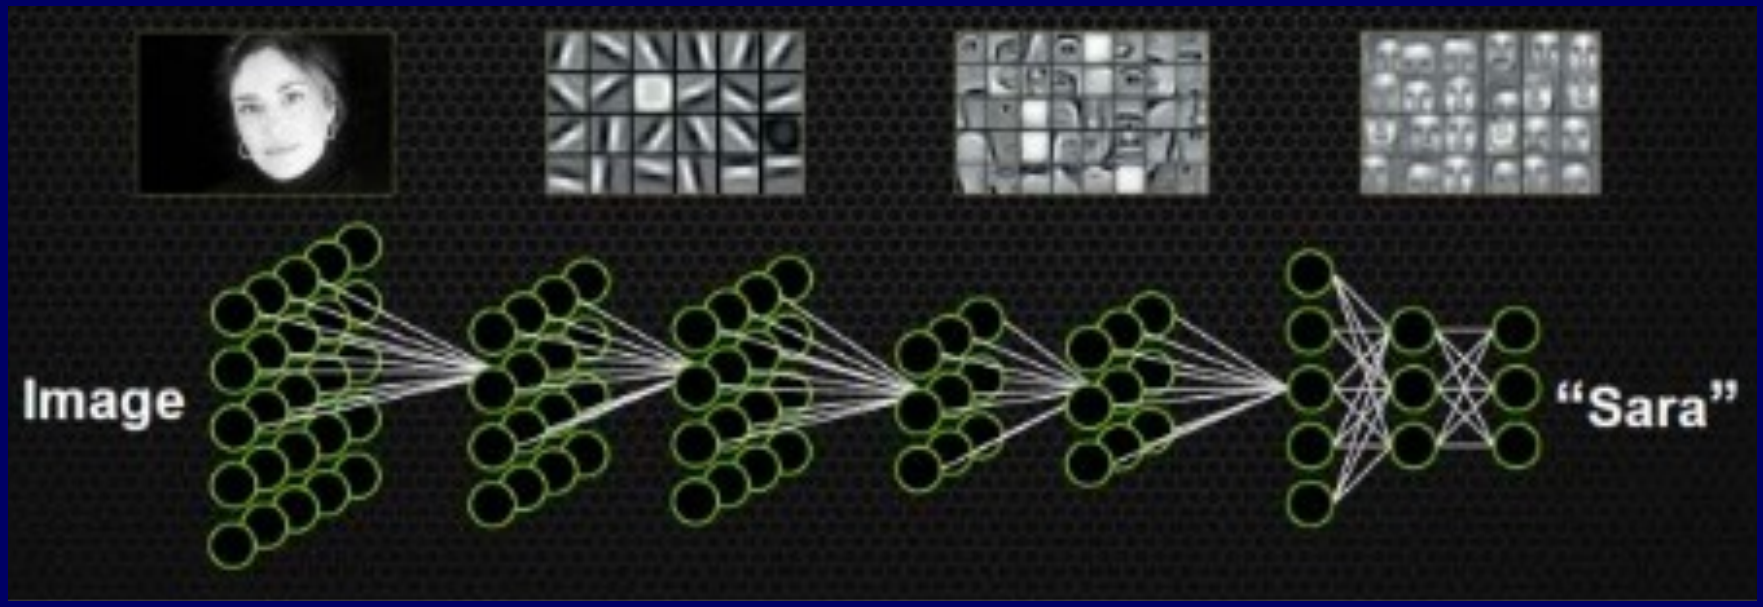

- Higher-level abstraction captured at higher layers
- Lower level features are combined at proceeding layers